



JavaPolis 2004

Enterprise JavaBeans 3.0

Linda DeMichiel

Sun Microsystems





Agenda

- Goals of EJB 3.0
- Overview of the Simplified EJB API
- Persistence Update
- Summary and Current Status





EJB 3.0 Goals

- Focus: Ease of Development
- Simplify EJB: make it easier to use
 - Simplified set of APIs
 - Reduce the number of classes the developer needs to produce
 - Eliminate deployment descriptor from developer's view
 - Improve developer productivity
- Capture broader range of developers
 - Make it simpler for average developer
 - Increase developer base, target more corporate developers





EJB 3.0 Goals

- Retain / improve power of EJB 2.1
- Ensure EJB 2.1 APIs are always available
 - Provide interoperability between new APIs and existing EJB 2.1 APIs
 - Provide for reuse of existing components in new applications
 - Support migration scenarios
 - Simplifications / enhancements to improve use of these APIs as well
- Entity Beans
 - Entity Beans are one of the more complex areas of EJB 2.1
 - Need to simplify / improve EJB persistence as well





Approaches to Simplification

- Leverage Java language metadata
- Configuration by exception
- Simplification of all bean types
 - Beans as “POJOs”
 - Elimination of unnecessary interfaces
- Simplification of environmental access
 - Injection
- Lighter-weight entities
 - Focus on object/relational mapping
 - Expansion of EJB QL
- More work done by container; less by developer





Metadata Annotations

- Java language metadata (JSR 175) introduced as part of J2SE™ 5
- Leveraged heavily in EJB 3.0
 - Eliminate need for developer to provide deployment descriptor
 - Demarcation of injection and interpositioning points for container behavior
 - Specification of object/relational mapping
 - Reduce the number of interfaces programmer needs to implement
 - Reduce the number of classes programmer needs to produce
 - Gives developer simple-to-use and yet powerful capability





Metadata Annotations

- Configuration by exception
 - Reduce need to explicitly specify common, expected behaviors and requirements on container
 - Leverage defaulting in conjunction with metadata
 - Defaults for annotation members
 - Metadata itself used for non-default cases (or for program documentation)
- Preservation of ability to use XML
 - As alternative mechanism
 - For overriding of annotations
 - For supplementing annotations





Simplification of EJB Bean Types

- More closely resemble “POJOs”
- Elimination of requirements for EJB component interfaces
 - Business interfaces are plain Java interfaces, not EJBObject/EJBLocalObject interfaces
- Elimination of requirement for home interfaces
 - Use dependency injection or simple bean lookup
- Elimination of requirement to implement SessionBean, MessageDrivenBean,... interfaces
 - Use annotations for (now optional) callbacks
- Elimination of all required interfaces for Entity Beans





EJB 2.1 Stateless Session Bean

```
public interface Calculator extends EJBObject {
    public int add(int a, int b)
        throws RemoteException;
    public int subtract(int a, int b)
        throws RemoteException;
}

public interface CalculatorHome extends EJBHome {
    public Calculator create()
        throws CreateException, RemoteException;
}
```





Stateless Session Bean Class

```
public class CalculatorBean implements SessionBean{
    private SessionContext ctx;
    public void setSessionContext(SessionContext ctx){
        this.ctx=ctx;
    }
    public void ejbCreate() {}
    public void ejbActivate() {}
    public void ejbPassivate() {}
    public void ejbRemove() {}
    public int add(int a, int b) {
        return a + b;
    }
    public int subtract(int a, int b) {
        return a - b;
    }
}
```





Deployment Descriptor

```
<session>
<ejb-name>CalculatorEJB</ejb-name>
<home>com.example.CalculatorHome</home>
<remote>com.example.Calculator</remote>
<ejb-class>com.example.CalculatorBean</ejb-class>
<session-type>Stateless</session-type>
<transaction-type>Container</transaction-type>
</session>
...
...
```





Same Example: EJB 3.0

```
@Stateless public class CalculatorBean  
    implements Calculator {
```

```
    public int add(int a, int b) {  
        return a + b;  
    }
```

```
    public int subtract(int a, int b) {  
        return a - b;  
    }
```

```
}
```

```
public interface Calculator {  
    public int add(int a, int b);  
    public int subtract(int a, int b);  
}
```





Same Example: Generated Interface

```
@Stateless public class CalculatorBean {  
  
    public int add(int a, int b) {  
        return a + b;  
    }  
    public int subtract(int a, int b) {  
        return a - b;  
    }  
}
```





Deployment Descriptor





Simplifications in this Example

- Elimination of Home interface
- Business interface is a plain Java interface
 - Bean can implement it or it can be generated
 - Bean can have more than one business interface
 - Can support remote access
 - EJB(Local)Object removed from client view
 - RemoteExceptions removed from client view
- No Implementation of SessionBean interface





Metadata and Defaulting

- Defaulting of transaction management types
- Defaulting of transaction attributes
- Default use of unchecked methods
- Default use of caller identity
- Defaulting of interface generation
- Defaulting of names
- Defaulting of O/R mapping (for Entities)
- Etc,...

Metadata annotations are available to further specify / customize all of these





Stateful Session Bean Example

```
@Stateful public class ShoppingCartBean {  
    private String customer  
  
    @Init public void startToShop (String customer) {  
        this.customer = customer;  
        ...  
    }  
  
    public void addToCart (Item item) {  
        ...  
    }  
  
    @Remove public void finishShopping () {  
        ...  
    }  
}
```





Simplified Environment Access

- Get JNDI APIs out of developer's view
- Express dependencies in metadata
- Injection of resources
 - Instance variable injection
 - Setter injection
- Simple lookup methods
 - lookup method added to EJBContext





Instance Variable Injection

```
@Stateless public class MySessionBean {  
    @Resource public DataSource customerDB;  
    public void myMethod(String myString) {  
        ...  
        Connection conn = customerDB.getConnection();  
        ...  
    }  
}
```





Setter Injection

```
@Stateless public class MySessionBean {  
    private DataSource customerDB;  
  
    @Resource  
    public void setCustomerDB(DataSource customerDB) {  
        this.customerDB = customerDB;  
    }  
  
    public void myMethod(String myString) {  
        ...  
        Connection conn = customerDB.getConnection();  
        ...  
    }  
}
```





Dynamic Lookup

```
@Inject
private void setSessionContext(SessionContext ctx) {
    this.ctx = ctx;
}
...
myShoppingCart = (ShoppingCart) ctx.lookup
    ("shoppingCart");
...
```





Entity Beans: Goals

- Simplify programming model
 - POJO/JavaBeans like other EJB 3.0 beans
- Support for light-weight domain modeling
 - Inheritance and polymorphism
 - Extension of query language
 - Focus on O/R mapping
- Make instances usable outside the container
 - Remove need for DTOs and similar anti-patterns
 - Facilitate testability
- Complete query capabilities





- Bulk update and delete operations
- Projection list (SELECT clause)
- Explicit joins (inner and outer)
- Group by, Having
- Subqueries (correlated and not)
- Additional SQL functions
 - UPPER, LOWER, TRIM, ...
- Dynamic queries

Support for native SQL queries in addition to EJB QL enhancements





Entity Beans

- Concrete classes (no longer abstract)
- Support use of new()
- No required bean interfaces
- getter/setter methods
 - Can contain logic, e.g., for validation, transformation, etc.
- Direct persistent instance variable access in bean
 - No exposure of instance variables outside bean class
- Collection interfaces for relationships
- Usable outside the container
- No required callback interfaces





EntityManager

- Serves as untyped “home”
- Similar in functionality to Hibernate Session, JDO PersistenceManager, etc.
- Methods for lifecycle operations
 - create, remove, merge, flush, etc.
- Factory for Query objects





Lifecycle Model

- **new()**
 - Instance is not yet managed
- **create()**
 - Instance becomes managed
 - Instance becomes persistent (in database) upon transaction commit
- **remove()**
 - Deletion of persistent instance upon transaction commit
- **Detached instances**
 - Instances exist outside of original transaction context in which created / retrieved
 - E.g., serialized to other tiers
- **merge()**
 - Merge detached instance state back into persistence context





Entity Bean Example

```
@Entity public class Customer {  
    private Long id;  
    private String name;  
    private Address address;  
    private HashSet orders = new HashSet();  
  
    @Id(generate=AUTO)  
    public Long getID() {  
        return id;  
    }  
  
    private void setID(Long id) {  
        this.id = id;  
    }  
}
```





Example (cont.)

```
public String getName() {
    return name;
}

public void setName(String name) {
    this.name = name;
}

@OneToMany(Cascade=ALL)
public Set<Order> getOrders() {
    return orders;
}

public void setOrders(Set<Order> orders) {
    this.orders = orders;
}
}
```





Client View

```
@Stateless public class OrderEntryBean {  
  
    @Inject private EntityManager em;  
  
    public void enterOrder(int cId, Order newOrder) {  
        Customer cust = em.find(Customer.class, cId);  
        cust.getOrders.add(newOrder);  
        newOrder.setCustomer(cust);  
    }  
  
    // other business methods  
}
```





Object/Relational Mapping Approach

What do we mean by this?

Persistence design spectrum:

Database centric approach



Object centric approach





O/R Mapping Approach

- Managed Java classes
 - provide an automated facility for developer working with domain object model mapped to relational database
- Developer is aware of mapping between database and domain object model
 - At query language level
 - At object/data transfer and faulting level
- Hooks are provided over mapping between relational database and Java object instances
 - Semantically explicit approach
 - Programmer is in control
 - Classes are transparent; mapping semantics are not





Application Developer's Strategy

- Definition of object model
 - In conjunction with database model to facilitate expected data access
- Construction and manipulation of “persistence context”
 - Set of managed persistent instances used by application within a particular scope
 - Application developer knows the needed shape of this
- Use queries to bring in data sets of interest
 - Most efficient way to access relational data
 - EJB QL significantly enhanced
 - NativeQuery methods also available for direct SQL queries
- Use metadata to control faulting patterns
 - Explicit control of data prefetch





Points of Control

- Over lifecycle
 - Cascade capabilities (CREATE, REMOVE, MERGE, ...)
- Scope of persistence context
- Queries
- Fetch / faulting behavior
 - FETCH JOINS
 - FetchType metadata
 - EAGER
 - LAZY
- Optimistic locking support
 - @Version; @Timestamp
- Isolation
 - FlushType, etc.





Example: Mapping Relationships

```
@Entity public class Employee {  
    private Address address;  
  
    @ManyToOne  
    public Address getAddress() {  
        return address;  
    }  
  
    public void setAddress(Address address) {  
        this.address = address;  
    }  
}  
  
@Entity public class Address {  
    ...  
}
```





Use of Defaulting

- Entity Employee is mapped to table EMPLOYEE
- Entity Address is mapped to table ADDRESS
- Table EMPLOYEE contains foreign key to table ADDRESS
- Foreign key column has same name and type as primary key of ADDRESS





Mapping Classes to Tables

- Use Java metadata to specify mapping
- Support for usual strategies
 - Table per class
 - Table per class hierarchy
 - Joined subclass
- Default type mappings available
- Custom type mappings under consideration





Example: Inheritance Mapping

```
@Entity
@Table(name="CUST")
@Inheritance(strategy=SINGLE_TABLE,
               discriminatorType=STRING,
               discriminatorValue="CUST")
public class Customer {...}

@Entity
@Inheritance(discriminatorValue="PCUST")
public class PreferredCustomer extends Customer {
    ...
}
```





Same Example, Using Defaults

@Entity

```
public class Customer {...}
```

@Entity

```
public class PreferredCustomer extends Customer {  
    ...  
}
```





Persistence Update: EJB 3 / JDO 2 Resolution

- Evolution of single persistence API based on EJB 3.0 draft
- Single persistence model, O/R mapping facility will be usable in J2EE and non-J2EE environments
- Persistence API in separate spec document, produced by JSR 220 (EJB 3.0) Expert Group
- Expansion of JSR-220 Expert Group
 - Inclusion of members from JDO Expert Group
 - Leading JDO vendors + individual members
 - Expert group is now ~35 members





Expansion of Scope: “Common API”

What It is:

- Single persistence model for common use in J2SE as well as J2EE
- Additional APIs to enable use outside of J2EE containers
 - APIs for acquiring EntityManager, managing transactions, configuration , etc

What It Isn't:

- It isn't “part” of J2SE (not in JDK)
- It doesn't shift primary focus from J2EE





Convergence Effort

What It Is:

- Merger of expertise
 - Leverage expertise of members from JDO expert community (as well as Hibernate, TopLink, EJB vendors,...)
 - Draw upon best of breed ideas from all these sources

What It Isn't:

- It isn't a merger of APIs

It is a merger of good ideas + experience





Current Focus

- Focus on O/R mapping
 - EJB 3.0 Early Draft as basis for expanded effort
 - Augment / improve / fill in holes
 - Fresh perspective of new members very helpful here
- Open Issues we're addressing
 - O/R mapping spec, including XML
 - o XML as alternative / overriding mechanism
 - SQL queries
 - APIs to enable use outside the container





EJB QL Enhancements

- Bulk update and delete operations
- Projection list (SELECT clause)
- Explicit joins (inner and outer)
- Group by, Having
- Subqueries (correlated and not)
- Additional SQL functions
- Dynamic queries





Examples: Bulk Update / Delete

```
DELETE  
FROM Customer c  
WHERE c.status = 'inactive'
```

```
UPDATE Customer c  
SET c.status = 'outstanding'  
WHERE c.balance < 1000
```





Examples: Subqueries

```
SELECT preferredCustomer
FROM Customer preferredCustomer
WHERE preferredCustomer.balance < (
    SELECT avg(c.balance) from Customer c)
```

```
SELECT DISTINCT emp
FROM Employee emp
WHERE EXISTS (
    SELECT spouseEmp
    FROM Employee spouseEmp
    WHERE spouseEmp = emp.spouse)
```





Examples: Joins

```
SELECT i.category
FROM Item i JOIN i.bids b
WHERE b.amount > :amount
```

```
SELECT i
FROM Item i LEFT JOIN FETCH i.bids
WHERE i.category = 'paintings'
```





Examples: Projection

```
SELECT c.id, c.status
FROM Customer c JOIN c.orders o
WHERE o.count > 100
```

```
SELECT new CustomerDetails(c.id, c.status, o.count)
FROM Customer c JOIN c.orders o
WHERE o.count > :ordercount
```





SQL Queries

- Allow direct SQL over actual database schema
 - Very useful for some applications
 - Database portability not important for some applications
 - Escape mechanism for very complex queries not handled by EJB QL
- Allow SQL query results to be mapped into entity beans and/or instances of other Java classes





Summary: EJB 2.1 Component View

- Contractual view of components
- Container specifies contracts
 - Interfaces beans must adhere to
 - Deployment descriptor contract
- Beans written to contracts
 - Whether they need them or not
- Very static viewpoint; negative consequences in terms of application complexity





Summary: EJB 3.0 Shift in Viewpoint

- Container is a service injection facility
 - Injection points specified through metadata
 - Container calls bean if/when bean specifies its need
- Inversion of contractual view
 - Bean doesn't implement a predefined set of APIs
 - Instead, bean specifies its expectations on the container
- More extensible architecture





Current Status

- Early Draft 1 released in summer
- Targeting Early Draft 2 in ~1 month
 - Additions:
 - **Callbacks**
 - **Interceptors**
 - **EntityManager infilling**
 - **Improvements to O/R mapping spec**
 - **SQL queries**
- Plan an Early Draft 3 as well
 - Use Early Drafts to get community feedback





Conclusion

- POJO-centric view of all enterprise beans
- Simplification of environment access
- Simplification of entity beans
 - Clear O/R mapping orientation
 - Improvements of query language capabilities
- Metadata is major enabling technology
- Major collaboration effort by Expert Group

