

Introduction to Clustering

At the end of this module, you will be able to:

- ✓ Understand the uses of a cluster
- ✓ Describe the basic cluster architecture
- ✓ Describe the multi-tier cluster architecture
- ✓ Understand how clusters communicate

Road Map



1. Cluster Architecture

- Cluster definition
- Basic cluster architecture
- Multi-tier cluster architecture
- Proxy servers

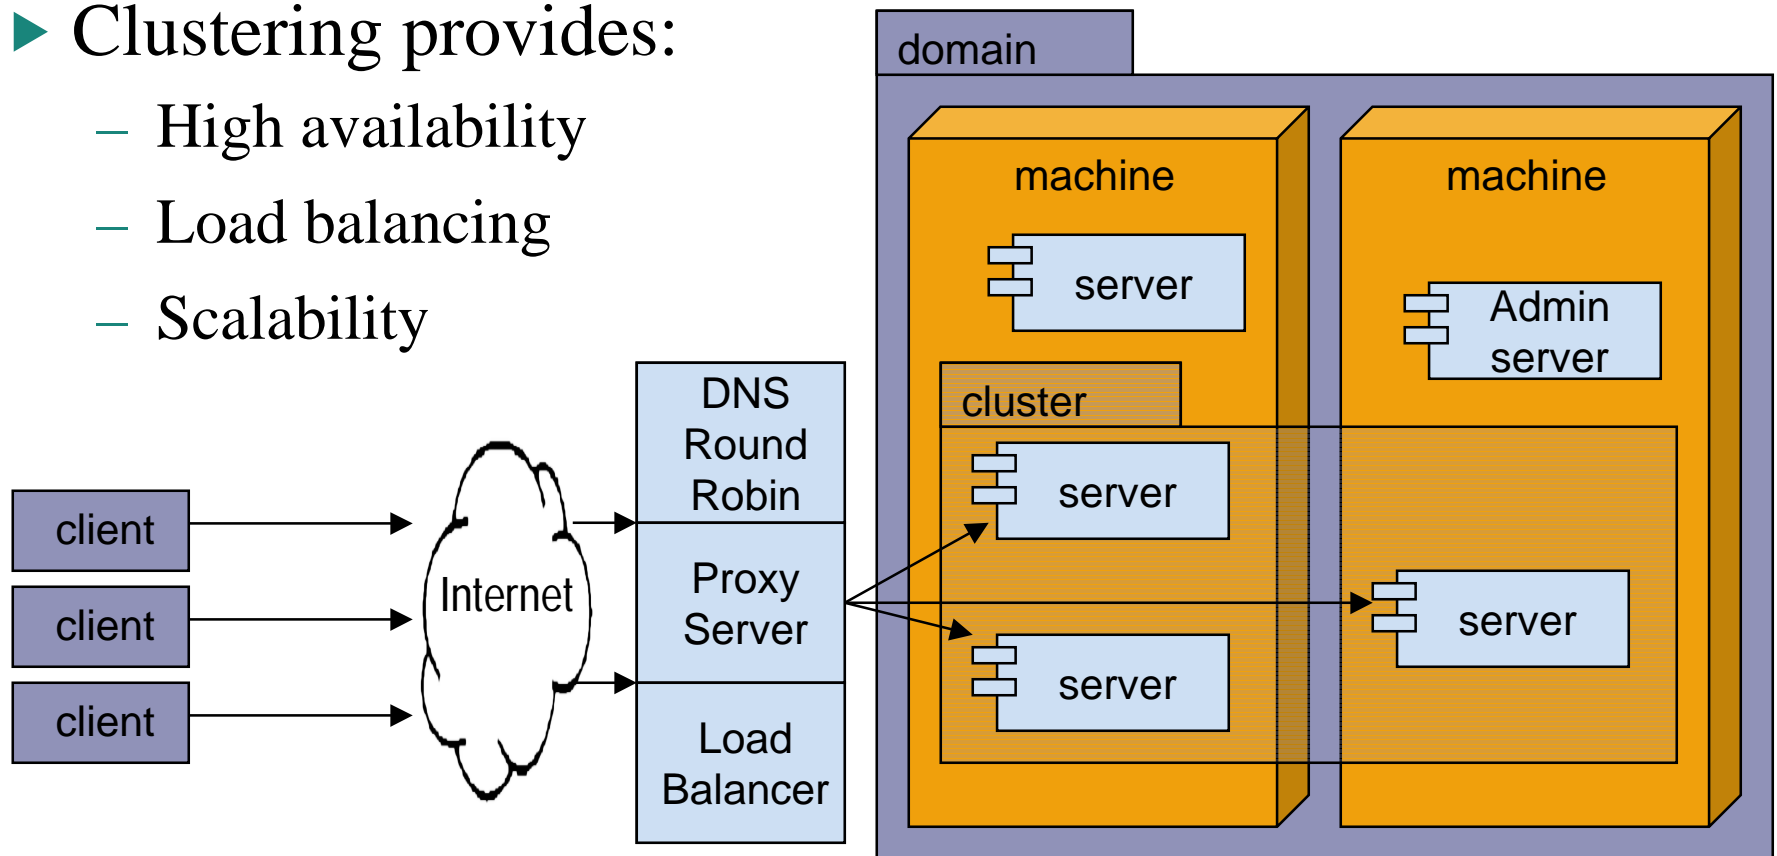
2. Networks and Clusters

3. Cluster Communication

Definition: Clustering



- ▶ A *cluster* is a group of WebLogic Server instances, working in coordination.
- ▶ Clustering provides:
 - High availability
 - Load balancing
 - Scalability



Benefits of Clustering



- ▶ There are two main benefits of clustering together WebLogic servers:
 - Scalability
 - High-Availability
- ▶ Scalability allows you to add and remove servers without disrupting current services.
- ▶ High-Availability ensures that when a server (in a cluster) fails, there are other servers to take over the work so the client is not affected.

Key Capabilities



- ▶ The key capabilities of a WebLogic cluster are:
 - Application Failover

When an object in an application is performing a task becomes unavailable, another object will take over and finish the job.
 - Site Failover

When all the services and applications in a single site fail they can switch to a separate site and continue processing.
 - Server Migration

When a server fails, pinned services can be migrated to another server in a cluster.
 - Load Balancing

The even distribution of tasks and communications across multiple servers.

- ▶ Applications are generally broken into multiple tiers, each representing their distinct functionality:
 - Web tier
 - Presentation tier
 - Business or object tier
- ▶ WebLogic provides clustering support for all three tiers.
- ▶ Other services, such as JMS and JDBC, can take advantage of clusters but load-balancing and failover is a little different.

Deciding on a Cluster Architecture




- ▶ Good architecture is somewhat subjective but there are a few global considerations:
 - Performance
 - Efficient replication
 - Optimal load balancing
 - Effective failover
 - Reliable communication
- ▶ There are two primary cluster architectures to choose from:
 - Basic Cluster architecture
 - Multi-tier architecture

Basic Cluster Architecture




- ▶ A basic cluster architecture combines static HTTP, presentation logic, business logic and objects into one cluster.



Multi-Tier Cluster Architecture



- The Web tier and the business logic with services can be separated into two clusters.



When to Use Multi-tier Architecture



- ▶ The multi-tier cluster is recommended for Web Applications that require:
 - Load balancing for method calls to clustered EJBs
 - Flexibility for load balancing between servers that provide HTTP content and servers that provide clustered objects
 - Higher availability (fewer single points of failure)
 - More flexible security

Basic Cluster Architecture Advantages and Disadvantages



- ▶ The basic cluster architecture has these advantages:
 - Easy administration
 - Flexible load balancing
 - Robust security

- ▶ The basic cluster architecture has these disadvantages:
 - Cannot load balance EJB method calls
 - Load balancing across the tiers may become unbalanced

Multi-tier Advantages and Disadvantages



- ▶ The multi-tier architecture has these advantages:
 - Improved load balancing
 - Load balancing of EJB methods
 - Higher availability
 - Improved security options

- ▶ The multi-tier architecture has these disadvantages:
 - Can create a bottleneck when presentation tier makes frequent calls to the business logic
 - Increased licensing cost
 - Added firewall configuration complexity

Proxy Servers




- ▶ Proxy servers are used to provide load balancing and failover for a cluster.
 - Are the client's first level of interaction with the cluster
 - Give the cluster its single server appearance
- ▶ A proxy server can be either software-based or hardware-based.
- ▶ A software-based proxy server may be an internal WebLogic servlet or a 3rd party application.
- ▶ A hardware-based proxy server is typically a physical load balancer.

Basic Cluster Proxy Architecture



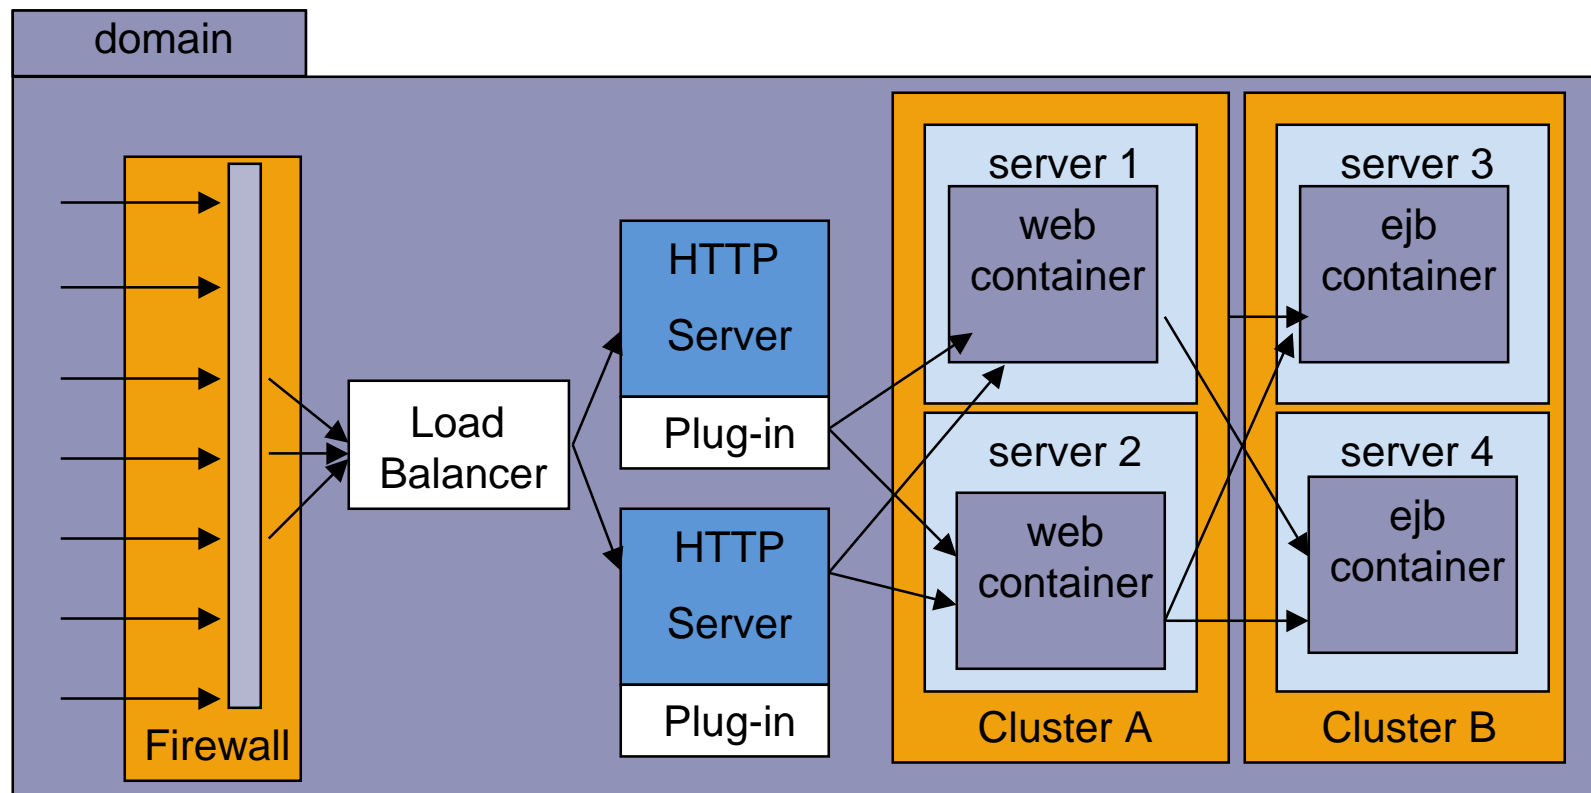
- ▶ Similar to the basic cluster architecture, except static content is hosted on non-clustered HTTP servers.



Multi-Tier Cluster Proxy Architecture



- ▶ Similar to the multi-tier cluster architecture, except static content is hosted on non-clustered HTTP servers.

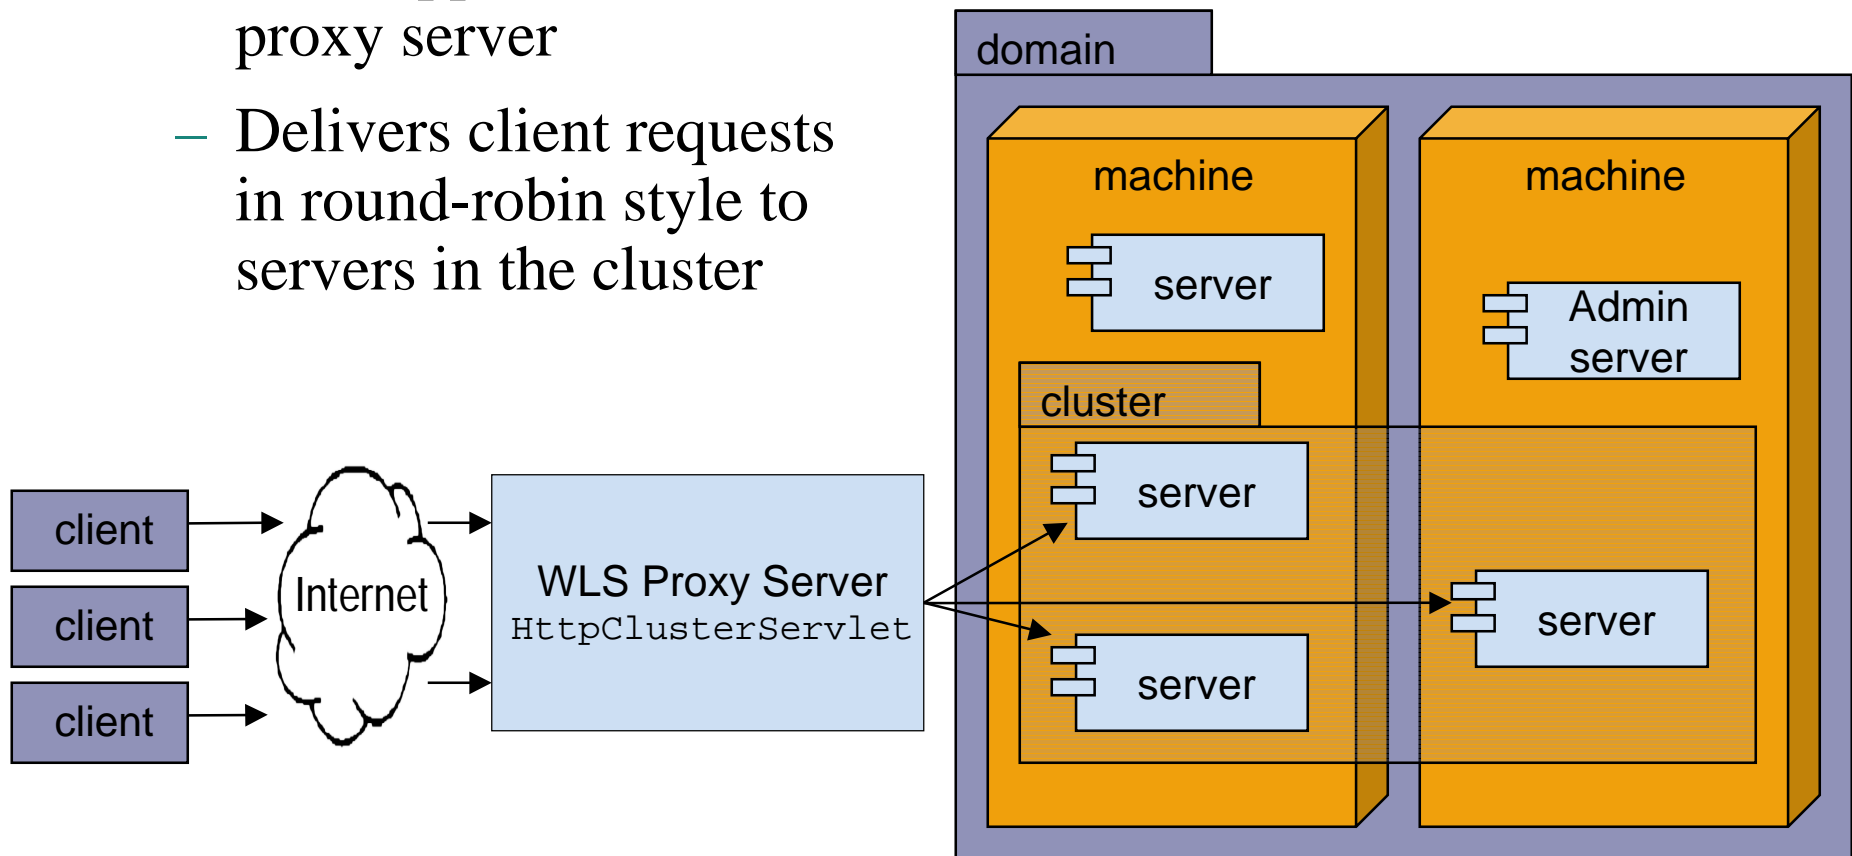


WLS HttpClusterServlet



► HttpClusterServlet:

- Is deployed in the default Web Application of the proxy server
- Delivers client requests in round-robin style to servers in the cluster



WLS Plug-Ins...



- ▶ WLS is compatible with major Web servers using the following plug-ins:
 - Sun Java System Web Server plug-in (formerly Netscape iPlanet or Sun One Web Server)
 - IIS plug-in (Microsoft IIS)
 - Apache plug-in

...WLS Plug-Ins



► Plug-ins:

- Delegate dynamic content requests to WLS
- Round-robin across a cluster
- Support routing based on URL path or on MIME type of the requested file or both
- Route HTTP requests to back-end WLS instances based on session cookie or URL rewriting
- Avoid failed servers in the cluster

Proxy Plug-in Vs. Load Balancer



- ▶ There are many advantages to using a physical load balancer instead of the proxy plug-in:
 - No need to configure client plug-ins
 - Eliminating the proxy layer reduces the number of connections
 - The availability of more sophisticated load balancing algorithms

- ▶ There are a number of disadvantages as well:
 - Additional administration
 - Explicit configuration of “sticky” sessions for stateful web applications

Architecture Recommendations



- ▶ If possible, place static web content on separate web servers in the DMZ.
- ▶ Use a combined tier architecture if your presentation and control tier makes multiple invocations of the business tier.
- ▶ Make sure that your architecture choice supports passing active and passive cookies between the cluster and client application.

Section Review



In this section, we learned how to:

- ✓ Describe the uses of a cluster
- ✓ Describe a basic cluster architecture
- ✓ Understand the use of a proxy server
- ✓ Use a multi-tier architecture
- ✓ Decide on the best cluster architecture to use



Road Map



1. Basic Clusters
2. Multi-Tier Clusters
3. **Networks and Clusters**
 - Local Area Networks
 - Metropolitan Area Networks
 - Wide Area Networks
4. Cluster Communication

Cluster in Networks




- ▶ WebLogic Server clusters can be created in three different kinds of networks:
 - Local Area Networks
 - Metropolitan Area Networks
 - Wide Area Networks
- ▶ When you are configuring your cluster, you will need to keep in mind the type of network you are using.

Local Area Networks




- ▶ A local area network (LAN) serves a local set of computers.
 - They usually use high quality, high-speed communication links
 - Typical data transmission speeds are 100 megabits/second
- ▶ Most clusters exist within a single LAN




Metropolitan Area Networks

- ▶ A Metropolitan Area Network (MAN) is a network that usually spans a campus or a city.
- ▶ You can have different clusters located reasonably close to each other within a MAN.



Wide Area Networks

- ▶ A Wide Area Network (WAN) usually spans a wider geographical area and can be made up of smaller MANs and Local Area Networks (LAN).
- ▶ You can have different clusters located in different regions within a WAN.
 - A cluster can be located in different LANs within a MAN or within a WAN.



Section Review



In this section, we learned how to:

- ✓ Describe different types of networks that a cluster can operate in



Road Map



1. Basic Clusters
2. Multi-Tier Clusters
3. MANs and WANs
- 4. Cluster Communication**
 - One-to-many communications
 - Peer-to-peer communications
 - Topic

Server Communication in a Cluster



- ▶ WebLogic Server instances in a cluster communicate with one another using two different techniques:
 - Multicast (UDP)
 - Sockets (peer-to-peer TCP)
- ▶ IP multicast broadcasts one-to-many communications among clustered instances.
- ▶ IP sockets are used for peer-to-peer communications between servers.

Detecting a Failure



- ▶ WebLogic clusters detect a failure of a server instance in the following ways:
 - Through the use of IP sockets
 - Through the WebLogic server heartbeat
- ▶ If a server in the cluster unexpectedly closes its socket, it will be marked as "failed" and its services will not be used.
- ▶ Server instances use multicast to broadcast heartbeats every 10 seconds to other server instances in the cluster.
 - If three heartbeats are missed from a peer server, the server is marked as "failed" and its services will not be used

One-to-Many Communications



- ▶ WebLogic Server uses one-to-many communication for:
 - Cluster-wide JNDI updates
 - Cluster “heartbeats”

- ▶ Because all one-to-many communications occur over IP multicast, when designing a cluster, consider these factors:
 - If your cluster spans multiple subnets, your network must be configured to reliably transmit messages
 - A firewall can break IP multicast transmissions
 - The multicast address should not be shared with other applications
 - Multicast storms may occur

Peer-to-Peer Communications



- ▶ WebLogic Server uses peer-to-peer communications for:
 - Accessing non-clustered objects that reside on a remote server instance in the cluster
 - Replicating HTTP session states and stateful session EJB states between a primary and a secondary server
 - Accessing clustered objects that reside on a remote server instance (typically, in a multi-tier cluster architecture)

Multi-Tier Communications



- ▶ Multi-tier clusters will require more IP Sockets than a combined-tier cluster:
 - One socket for replicating session states
 - One socket for each WebLogic Server in the EJB cluster, for accessing remote objects

- ▶ As an example, using a three-node cluster, the worst-case scenario would be five open sockets per server:
 - One primary and secondary replicated session
 - Each server simultaneously invokes a remote EJB method on each node in the cluster

Communication in a WAN



- ▶ In a WAN, the servers in your cluster may span multiple subnets.
- ▶ In order for multicast messages to reliably transmit across the WAN your network must meet the following requirements:
 - Full support of IP multicast packet propagation
 - A network latency that allows for multicast messages to reach their destination in 200 to 300 milliseconds
 - A multicast time-to-live value high enough to ensure that routers do not discard multicast packets

Section Review



In this section, we learned how to:

- ✓ Describe how servers communicate in a cluster
- ✓ Understand how multi-tier clusters communicate



Module Review



In this module, we learned how to:

- ✓ Describe the differences between a basic cluster architecture and a multi-tier cluster architecture
- ✓ Use a proxy server to load balance requests to a cluster
- ✓ Describe how servers in a cluster communicate

