Clustering Services

At the end of this module, you will be able to:

- ✓ Set up whole server migration
- ✓ Understand JNDI clustering
- ✓ Understand JDBC clustering
- ✓ Migrate a JMS server
- ✓ Migrate transactions

Road Map

1. **Clustering Services**
   - Clusterable Services
   - Service Level Migration
   - Server Level Migration
2. JNDI
3. JDBC
4. Transactions
5. JMS

Which Services Can Be Clustered?

- A *clustered service* is an API or interface that is available on multiple servers in a cluster.
- WebLogic Server provides clustering services for:
  - Web applications
  - EJB and RMI objects
  - JNDI tree
- WebLogic Server provides limited clustering services for:
  - JDBC connections
  - JMS connection factories

Service Failover

- Services that are deployed to all servers in a cluster automatically have transparent failover.
- JMS and JTA cannot be deployed to all servers in a cluster
  - They are pinned services
- There are two ways to handle JMS and JTA failover:
  - Service Level Migration
  - Server Level Migration

Service Level Migration

- Service-level migration is performed manually.
- Pinned services are moved from the failed server to a new server.
  - The new server then takes over and processes any data left over when the server failed then takes over any future processing
- A migratable service is accessed by clients using a migration-aware RMI stub.

Server Level Migration

- Server level migration moves the whole clustered server instance, in it's entirety.
  - It is an alternative to service level migration
  - They should not be used together
- Server level migration is:
  - Only supported using the Secure Shell (SSH) version of Node Manager
  - Not supported on Windows
- Server level migration can happen both manually and automatically.
- When a migratable server fails:
  - The server is automatically restarted on the same machine
  - If it cannot be started on the same machine, it moves to another machine

Server Level Migration Services and Resources

- Migratable Servers
  - Managed servers that host pinned services
- Cluster Master
  - One server instance that orchestrates the process of automatic server migration
- Target Machines
  - Allowable or preferred hosts for migratable servers
- Node Manager
  - A process that starts and stops server instances
- Lease Table
  - A database table where migratable servers persist their state information
- Administration Server
  - Configures migratable servers and target machines as well as obtains the state of migratable servers and orchestrates the manual migration process

Server Migration Environment

- Before configuring server migration, your environment must be set up properly.
  - Each managed server must use the same subnet mask for multicast communications
  - Each server hosting migratable services should be time synched
  - All operating system versions must support identical functionality for `ifconfig`
  - Interface names for each server should be the same
  - The database you are using to support migration must be supported

Server Migration Configuration…

- To configure server migration:
  - Configure each machine using server migration with a floating IP address
  - Configure the SSH version of node manager on each managed server (it must be running as well)
  - Configure the database by creating the leasing tables used to store machine-server associations
  - Configure a data source to point to the server migration database
  - Grant superuser privileges to the `wlsifconfig.sh` script
  - Ensure that `wlsifconfig.sh, wlscontrol.sh and nodemanager.domains` are included in your machine's PATH

…Server Migration Configuration…

- To configure server migration (continued):
  - Ensure that the machines hosting migratable services trust each other
  - Set the candidate machines for server migration.
  - Edit `wlscontrol.sh` and set the Interface variable to the name of your network interface
  - Configure the server properties to enable server migration

…Server Migration Configuration

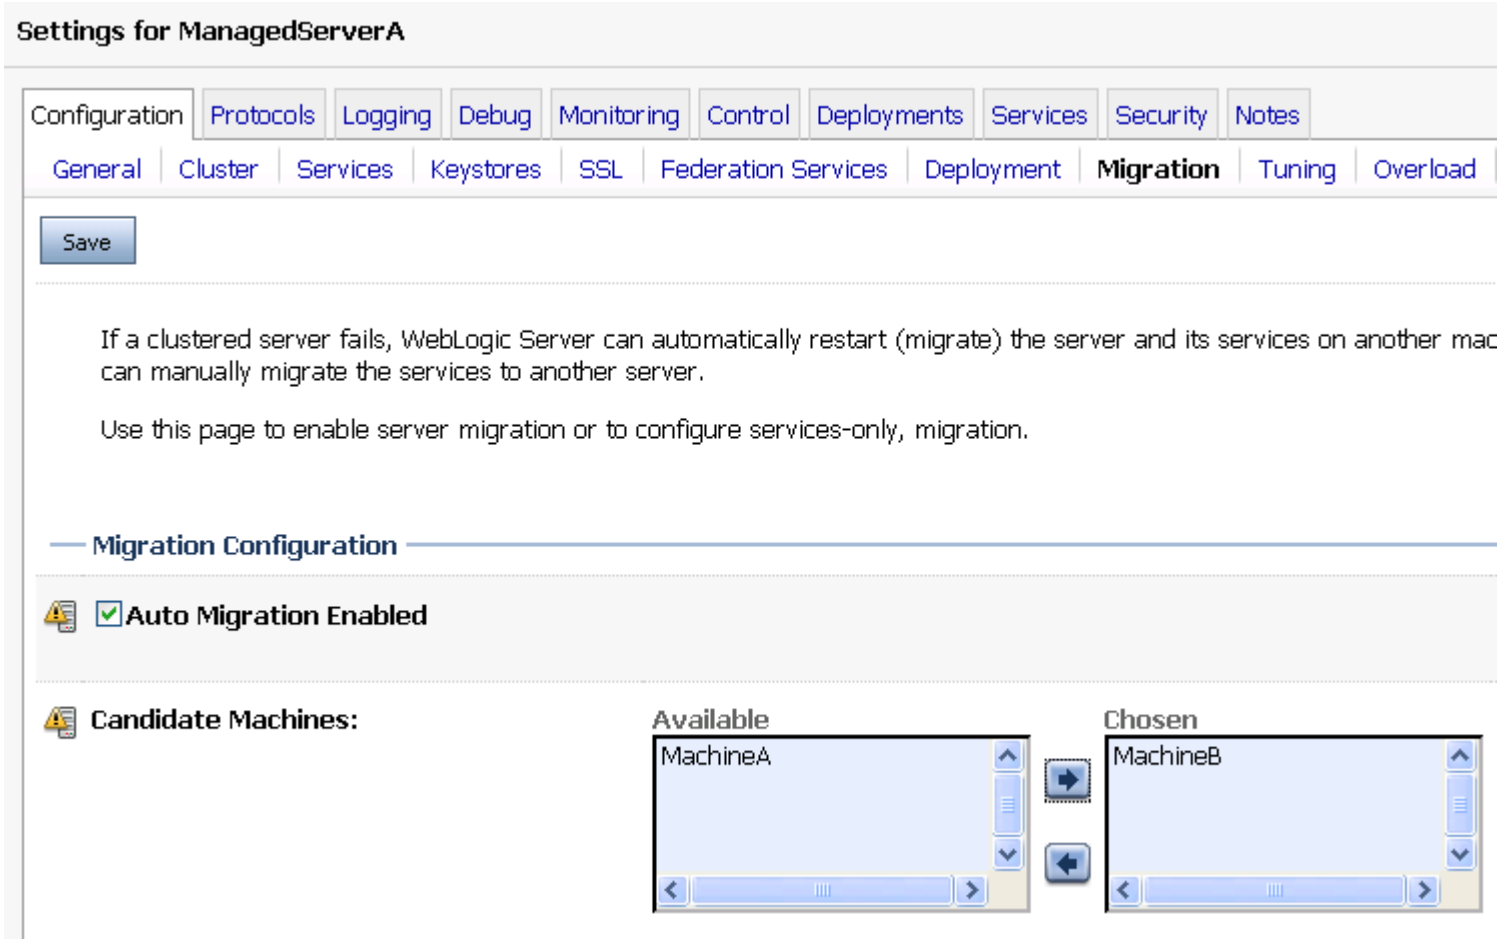

Server Migration Restrictions

- Each managed server must be configured to use the same subnet mask.
- Channels/Network Access Points must not have different Listen Addresses on a migratable server
- There is no built-in mechanism for transferring files
  - It is recommend that you store files on a different disk that is accessible by all the machines
- Node Manager security files must be copied to each machine
  - You can use the `nmEnroll()` WLST command

Migratable Server Startup Process

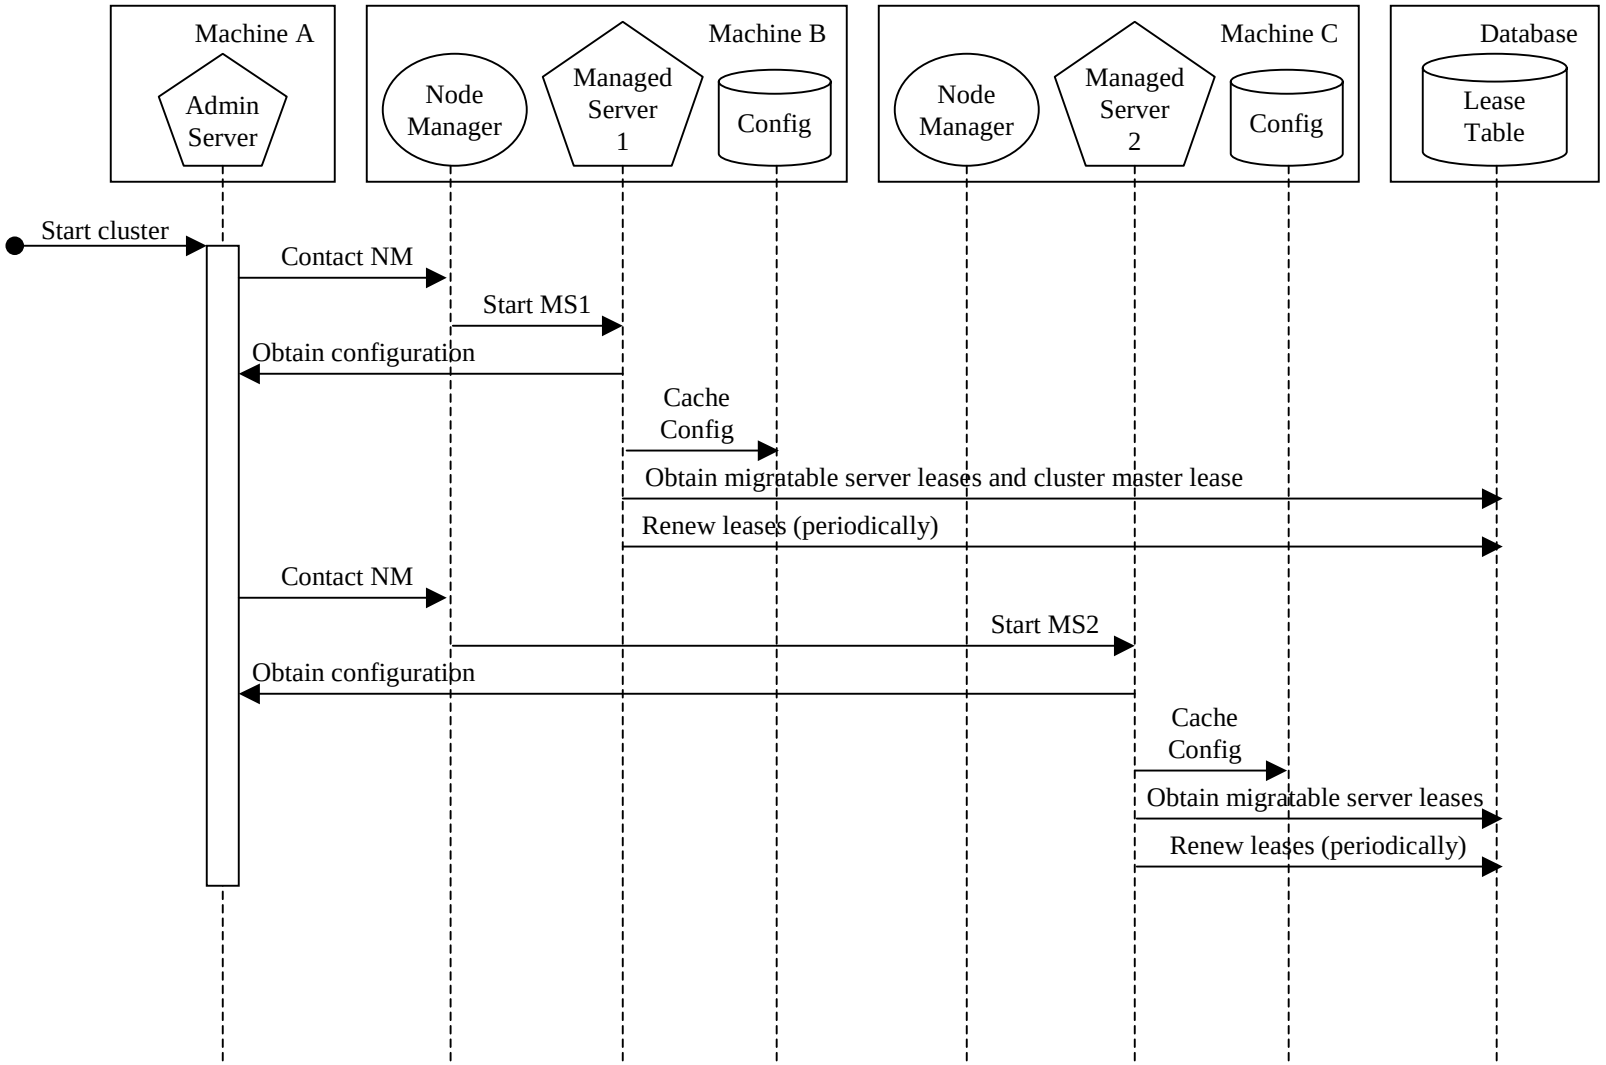

Automatic Server Migration

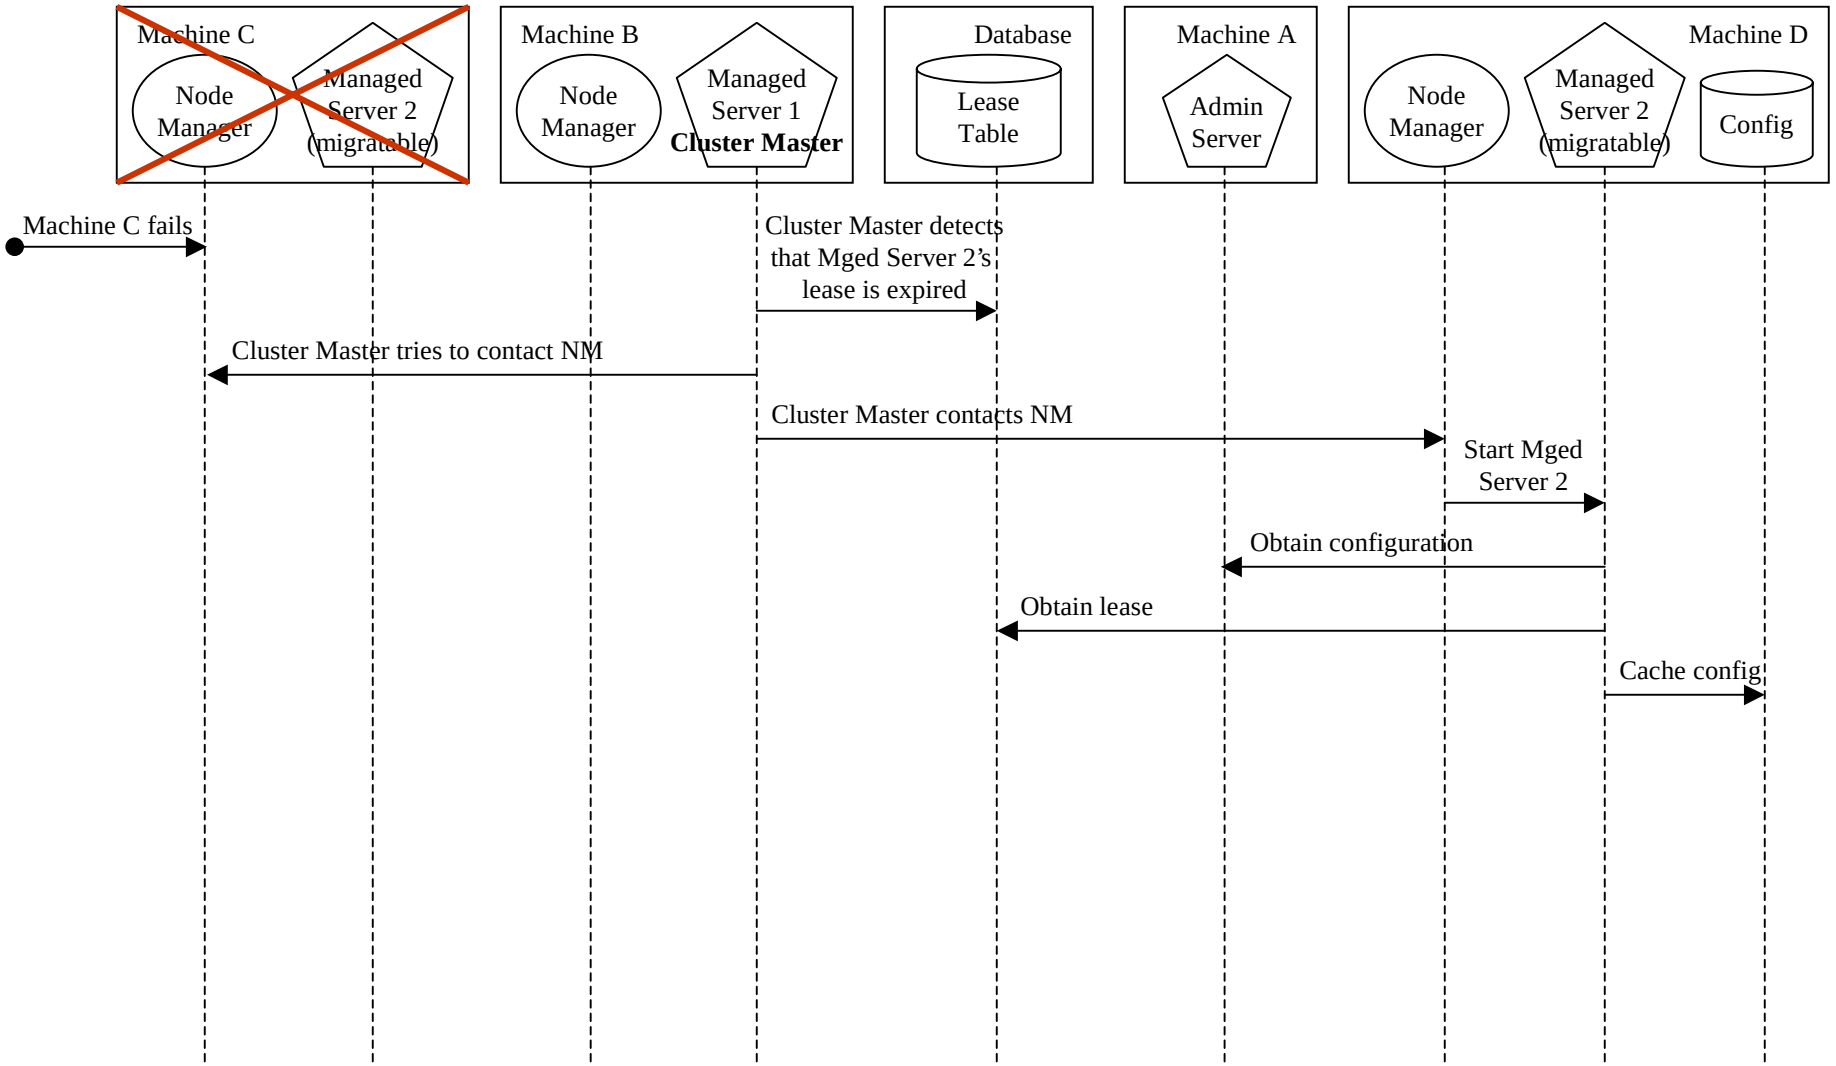

Manual Server Migration

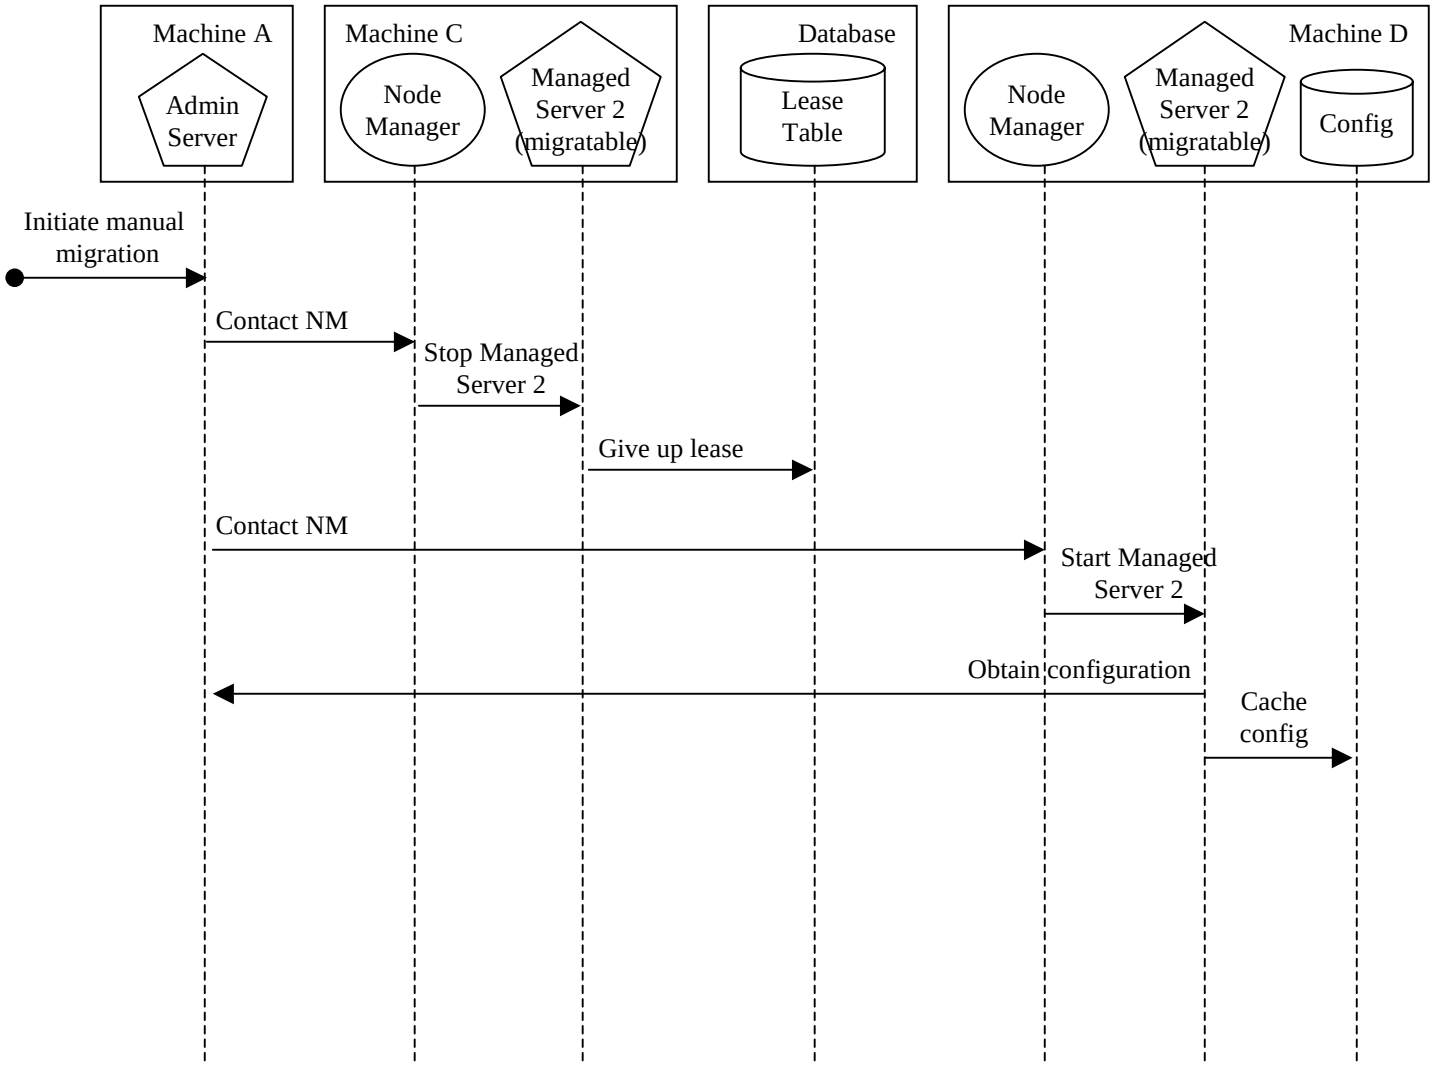

Section Review

In this section, we learned how to:

- ✓ Understand what services are clusterable
- ✓ Understand the difference between service level migration and server level migration
- ✓ Set up server level migration

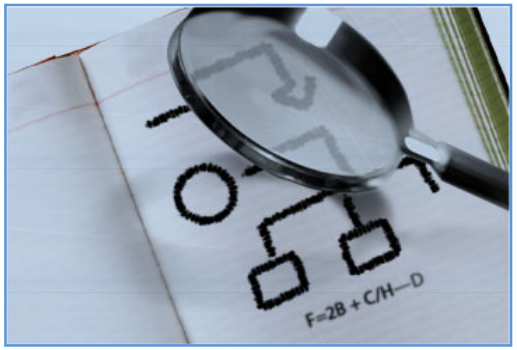

Road Map

1. Clustering Services
2. **JNDI**
   - Cluster-Wide JNDI Server
   - JNDI Naming Conflicts
3. JDBC
4. Transactions
5. JMS

JNDI Clustering Support

- The JNDI tree contains factory objects for accessing these items:
  - JDBC DataSources
  - EJBs, RMI objects
  - JMS connection factories
- WebLogic Server replicates clusterable objects to all servers on a cluster, transparently.

Cluster-Wide JNDI Service

- Each WLS server issues messages that announce these events:
  - A new object bound into the naming tree
  - An object removed from the naming tree
  - An object rebound (updated) into the naming tree

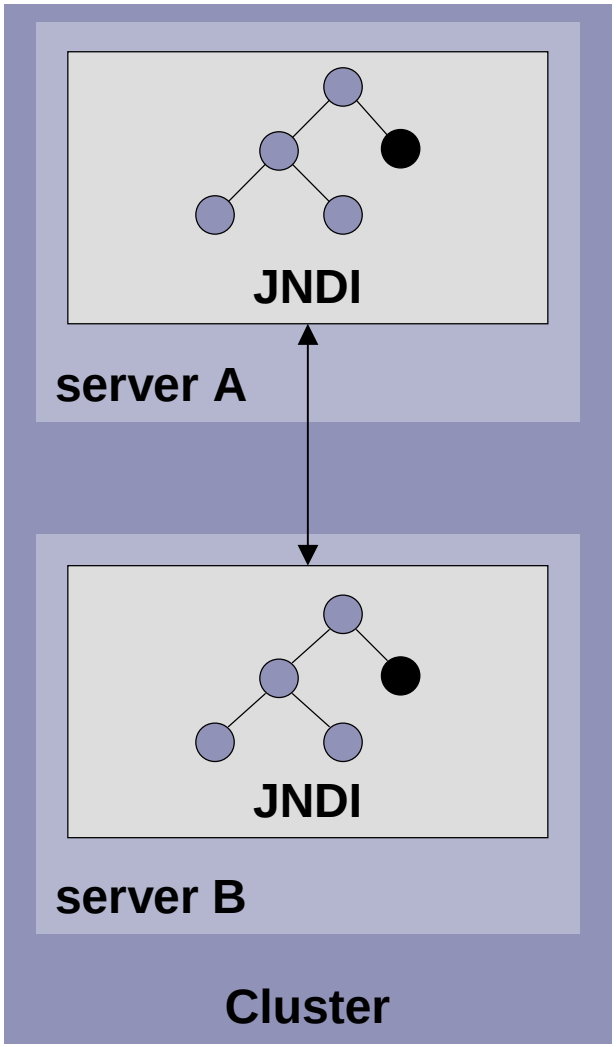

JNDI Naming Conflicts

- A naming conflict generally occurs when a server instance attempts bind a service to the JNDI when that name already exists.
  - The services are non-clustered
- In a cluster, a naming conflict can occur when the server tries to bind a clustered object with the same name as a non-clustered object.
  - The object will be bound to the JNDI locally
  - Other servers will not bind the replica-aware stub
- To avoid naming conflicts:
  - Deploy all clusterable objects to all servers in the cluster
  - Deploy to the cluster itself, not the individual servers

Section Review

In this section, we learned how to:

- ✓ Understand how JNDI clustering works
- ✓ Deal with JNDI naming conflicts

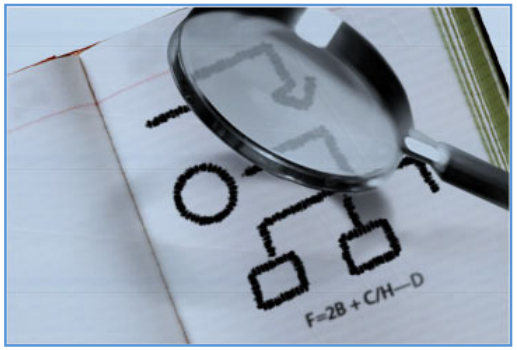

Road Map

1. Clustering Services
2. JNDI
3. **JDBC**
   - JDBC Clustering
   - Clustering vs. Multi-Data Sources
   - Targeting a Data Source to a cluster
4. Transactions
5. JMS

JDBC Clustering

- These JDBC objects are clusterable:
  - Data Sources
  - Multi Data Sources
- When you target a JDBC data source, a new instance of the data source is created on the target.
  - Single server target—a data source instance is created on the server.
  - Cluster target—a data source instance is created on all member servers in the cluster.
- Clustering your JDBC objects does not enable failover of connections but it can ease the process of reconnection when a connection fails.

Clustering vs. Multi Data Sources

Feature	Cluster	Multi Data Source
Local Clients	Access via JNDI	Access via JNDI
Remote Clients	Access via JNDI	Access via JNDI
Failover	None	On the same server (*High-Availability Algorithm only)*
Load balance	None	On the same server *(Load-Balance Algorithm only)*

Targeting a DataSource to a Cluster...

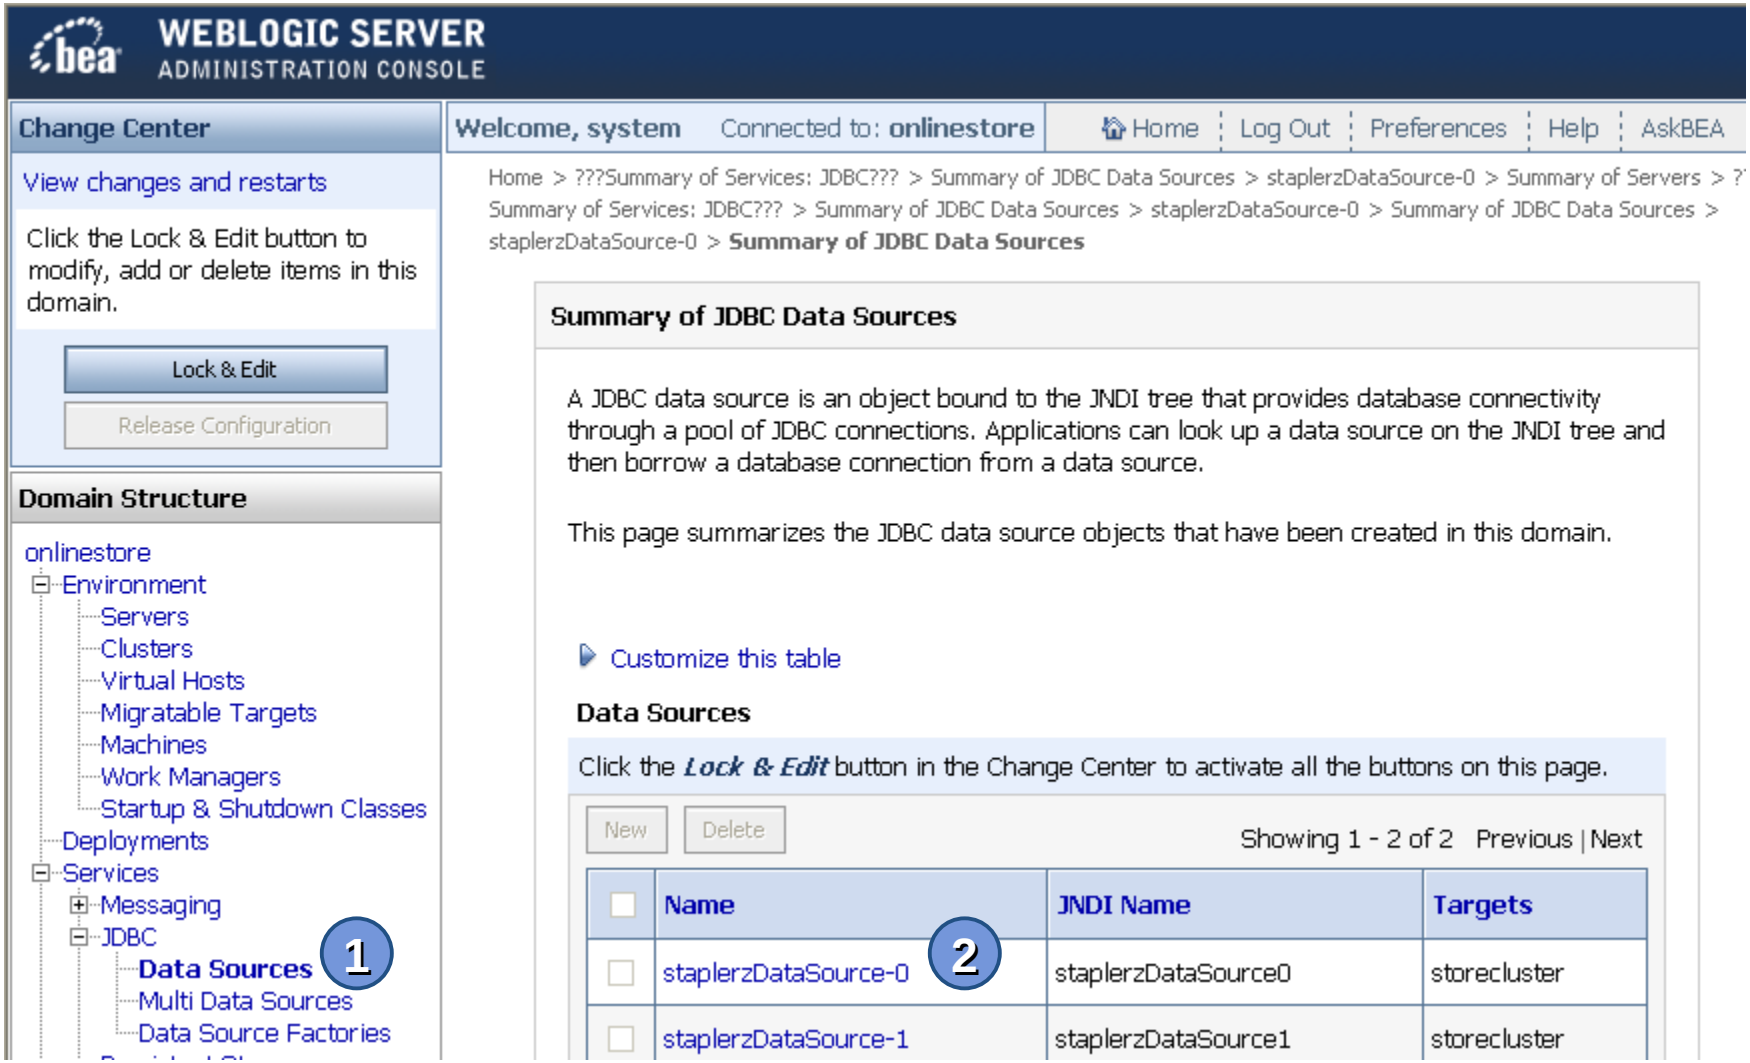

...Targeting a DataSource to a Cluster

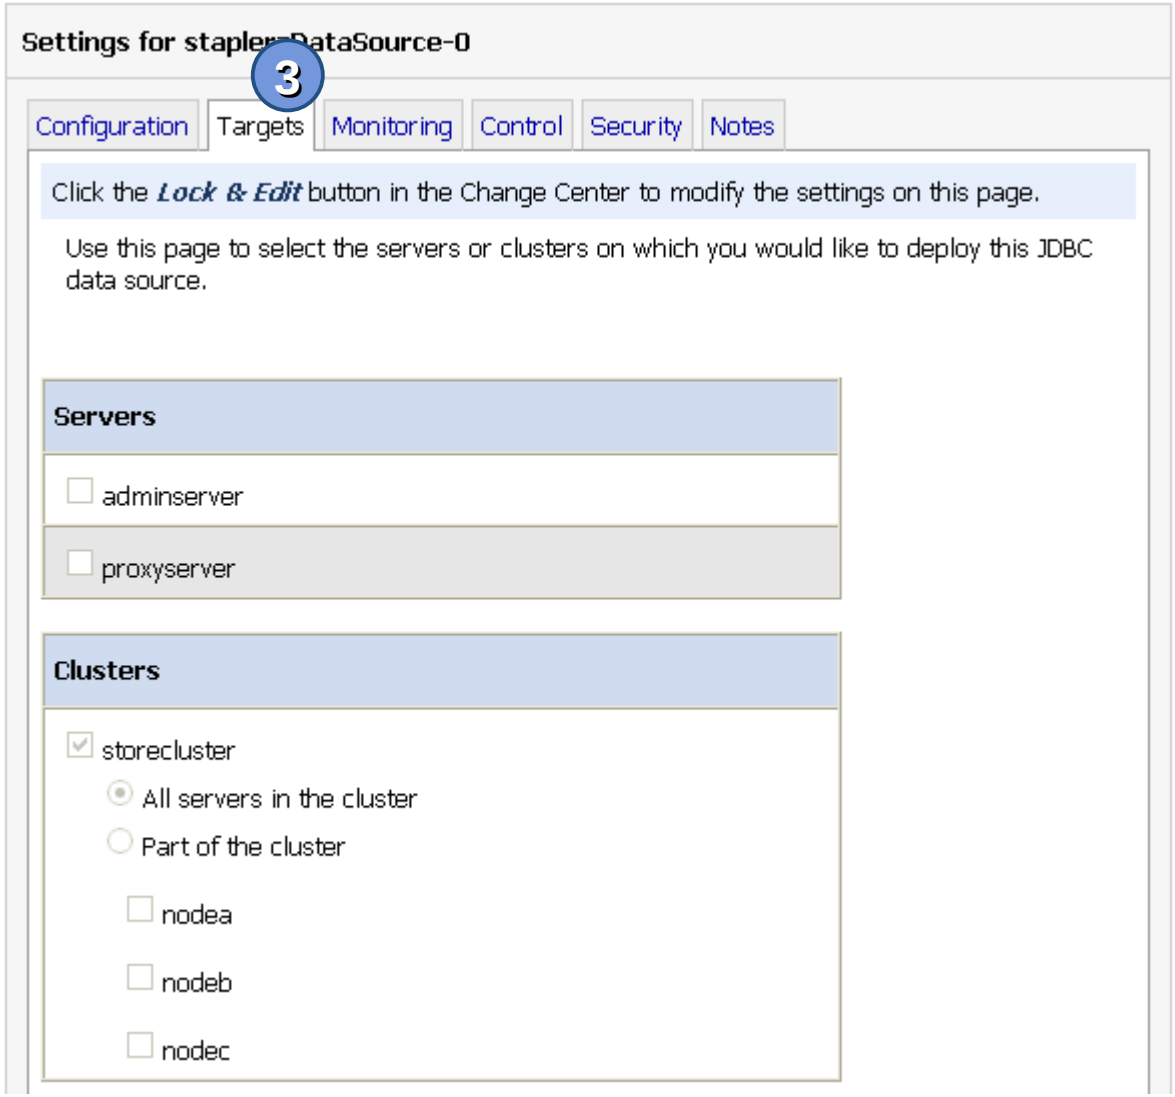

Section Review

In this section, we learned how to:

- ✓ Understand JDBC clustering capabilities
- ✓ Understand the difference between JDBC clustering and multi data sources
- ✓ Deploy a JDBC data source to a cluster

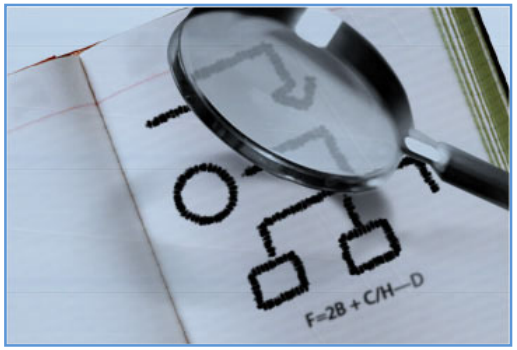

Exercise

Cluster JDBC Data Sources

- For details on the exercise, refer to the Lab Guide.
- If questions arise, ask the instructor.
- The instructor will determine the stop time.

Road Map

1. Clustering Services
2. JNDI
3. JDBC
4. **Transactions**
   - Transaction Recovery Service
   - Migrating JTA
5. JMS

Transaction Recovering After Failure

- The Transaction Manager (TM) makes every effort to resolve prepared transactions even after multiple crashes.
- The TM uses the Transaction Recovery Service (TRS) which:
  - automatically attempts to recover transactions on system startup
  - owns the transaction log for a server
  - parses all log files on startup for incomplete transactions and completes them

Transaction Recovery Service

- When a server restarts after a crash or if the JTA is migrated, the TRS attempts to:
  - resolved prepared transactions in T-Logs
  - maintain consistency across resources
  - persist in achieving transaction resolution
  - report heuristic completions

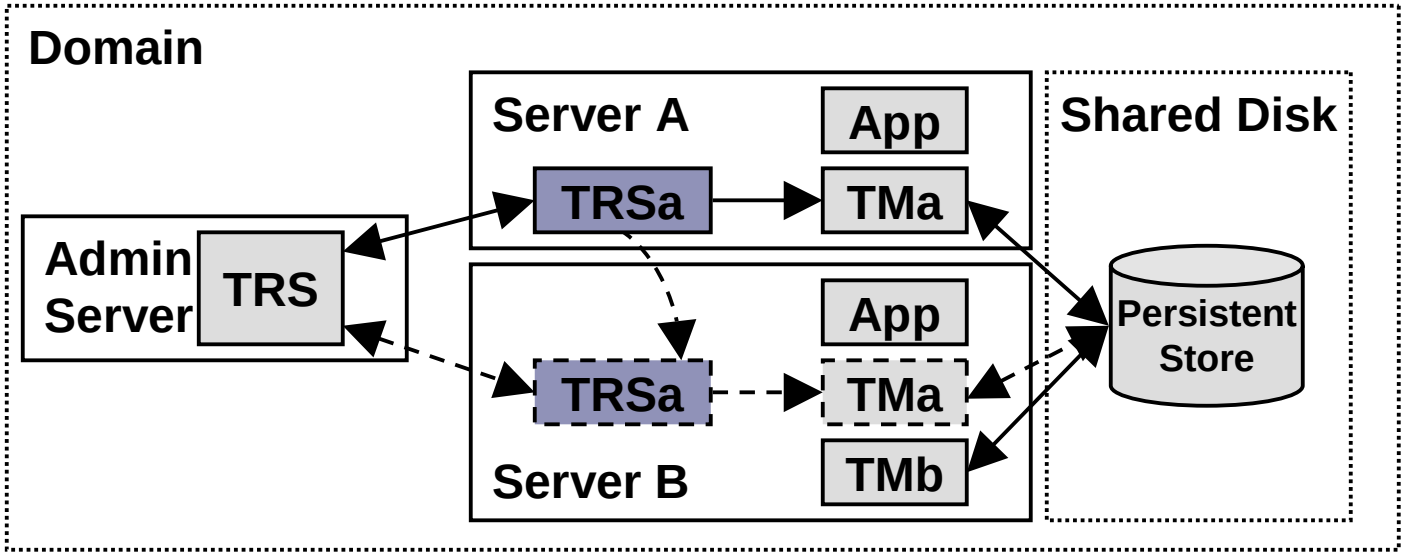

Recovering JTA Without a Cluster

- To recover transactions from a non-clustered server:
  - Make the persistent store for the failed server available to the new server
  - Set the path for the default persistent store to the path to the data file
  - Start the new server
- The TRS searches all transaction log files for incomplete transactions.

Recovering Transactions in a Cluster

- When recovering transaction in a clustered server you have two options:
  - Transaction Recovery Service migration
    - the TRS on backup server takes ownership of transaction log from crashed server
    - the TRS searches all transaction log files from failed server and attempts to complete in-flight transactions
    - if TRS on the backup server successfully completes all incomplete transactions, the server releases ownership of the TRS
  - Whole Server Migration
    - migrates the server in its entirety, along with all the services it hosts

Recovering Transactions for a Failed Clustered Server (Manually)

- Manually migrate the TRS from the crashed server to another server in the cluster using the Admin Console or the command-line interface:
  1. The TRS on the backup server takes ownership of the transaction log from the failed server.
  2. The TRS searches transaction log files from the failed server for incomplete transactions and completes them.
  3. If the TRS on the backup server successfully completes all incomplete transactions from the failed server, it releases ownership of the TRS (including transaction log files) for the failed server, so the failed server can reclaim it upon restart.

Configuring JTA Service Migration

- Before you can manually migrate JTA, you can to define which servers are available for migration.
- To set up the candidate servers use:

  Environment → Servers → Configuration → Migration

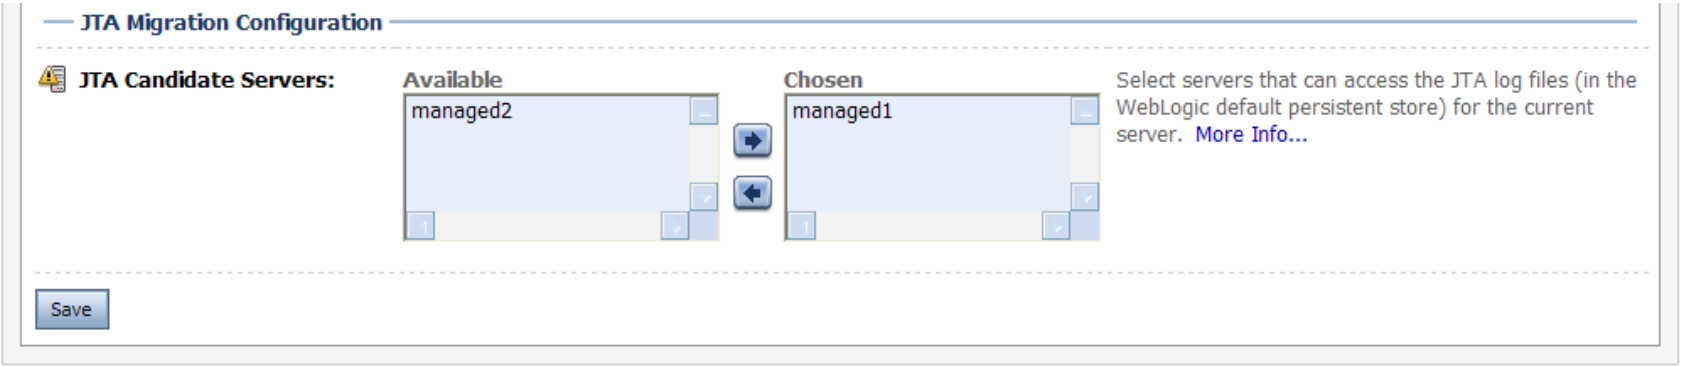

Manually Migrating the JTA to Another Server in a Cluster

- When you need to migrate transactions from a failed server to a working server use:

  Environment → Servers → Control → Migration

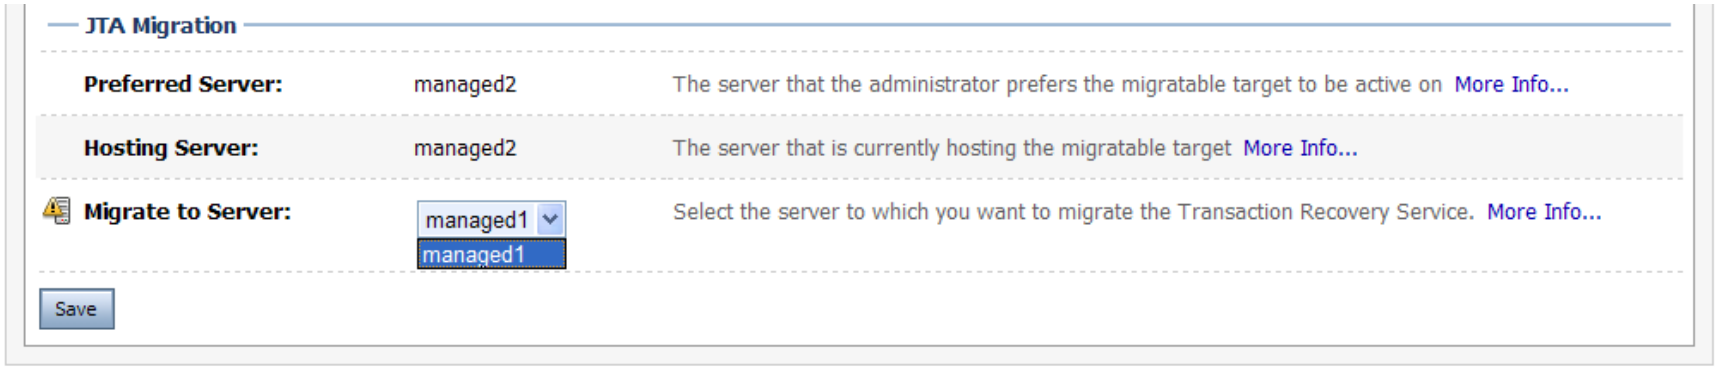

JTA Migration Limitations

- Migrating the JTA has the following limitations:
  - The JTA cannot be migrated to a backup server from a running server
    - The server must be stopped before migrating the JTA
  - The backup server only processes incomplete transactions
    - The backup server does not accept new transaction work for the failed server
  - The backup server does not process heuristic log files

Section Review

In this section, we learned how to:

- ✓ Understand how the transaction recovery service recovers transactions after a failure
- ✓ Manually migrate transactions to another server in the cluster

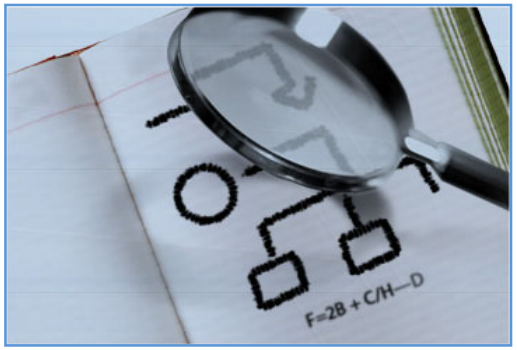

Road Map

1. Clustering Services
2. JNDI
3. JDBC
4. Transactions
5. **JMS**
   - Clustering JMS
   - Distributed Destinations
   - JMS Failover

JMS Clustering Support

- WebLogic Server supports targeting JMS connection factories to a cluster.
- WebLogic Server supports distributing JMS destinations throughout a cluster (Distributed Destinations)
  - JMS Queues and Topics are still managed by a single server instance in the cluster.
- WebLogic Server provides fail-over for JMS messages through:
  - Distributed destinations
  - Server-level automatic migration

JMS Connection Factory Clustering...

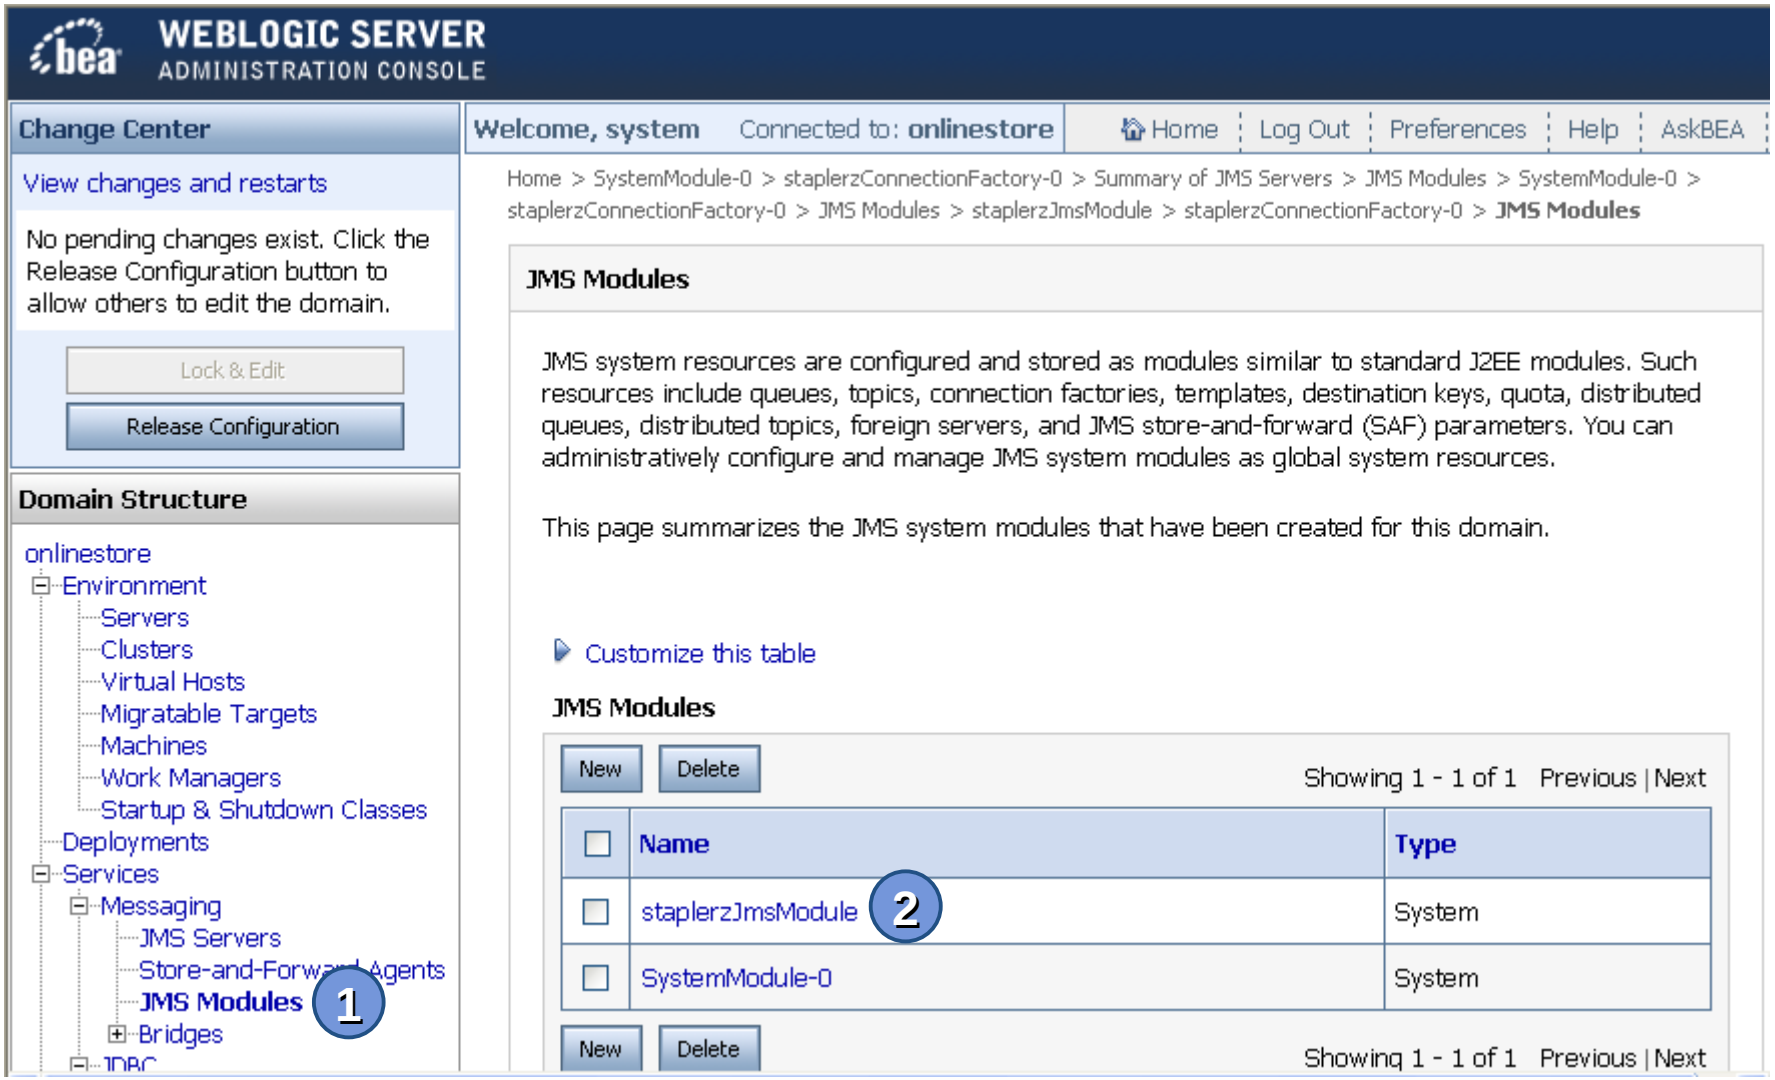

...JMS Connection Factory Clustering

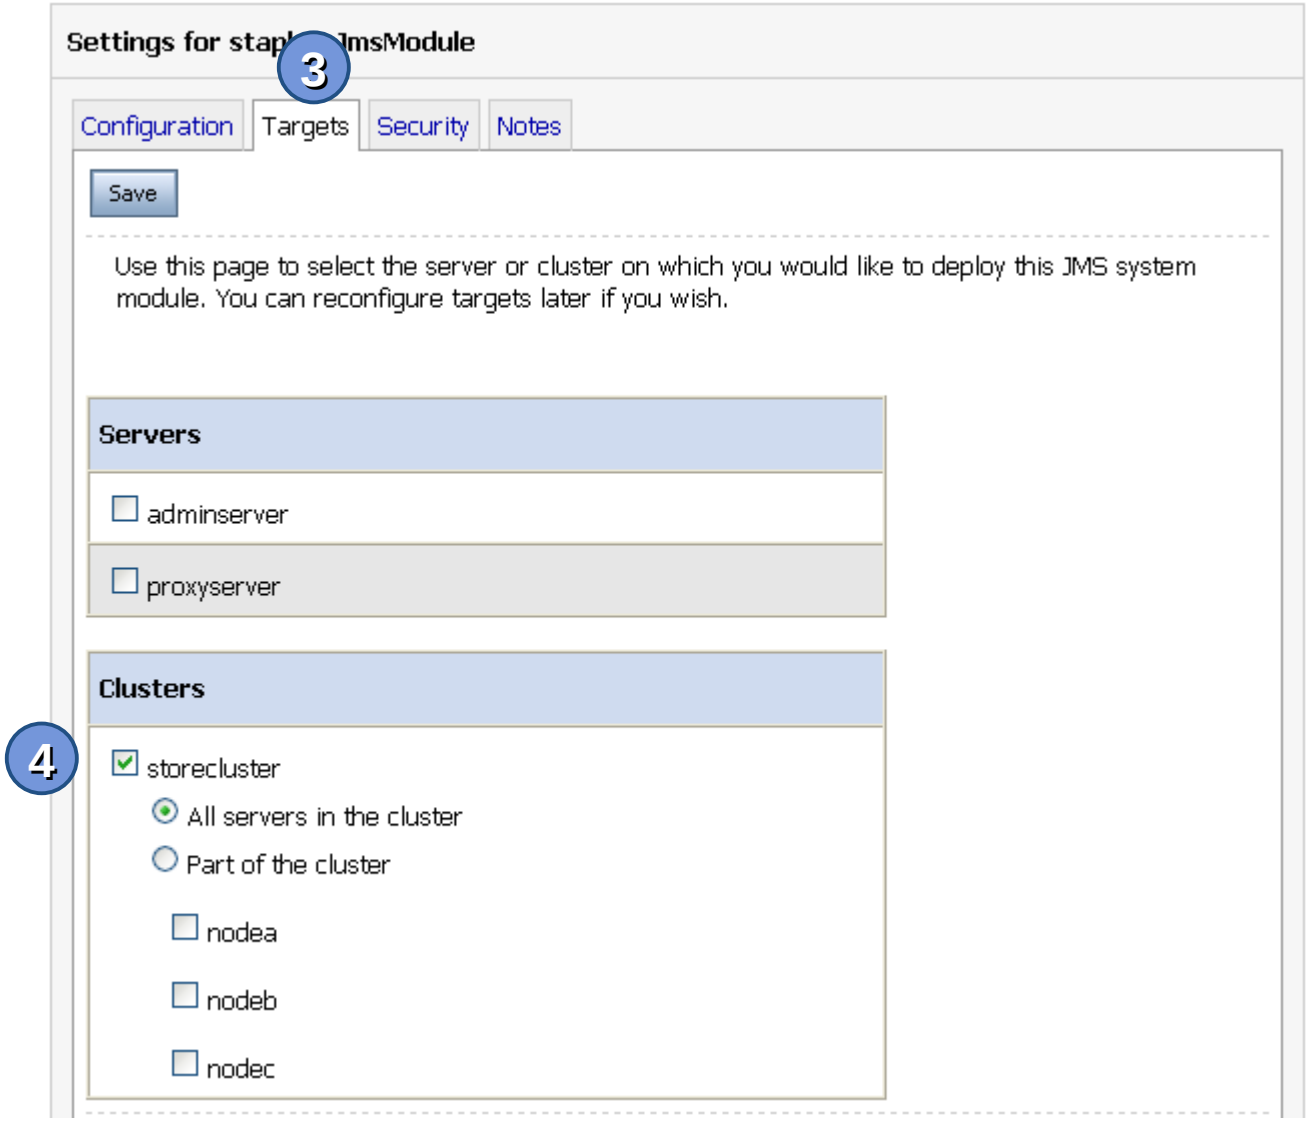

Distributed Destination

- A distributed destination has these properties:
  - It defines multiple destinations as part of one distributed destination.
  - It is looked up as a regular destination, via JNDI.
  - Its member availability is dynamically updated.
- Producers and consumers can send to and receive from a distributed destination.
- WebLogic JMS distributes the load across the available physical destinations within the distributed destination.

Distributed Queues

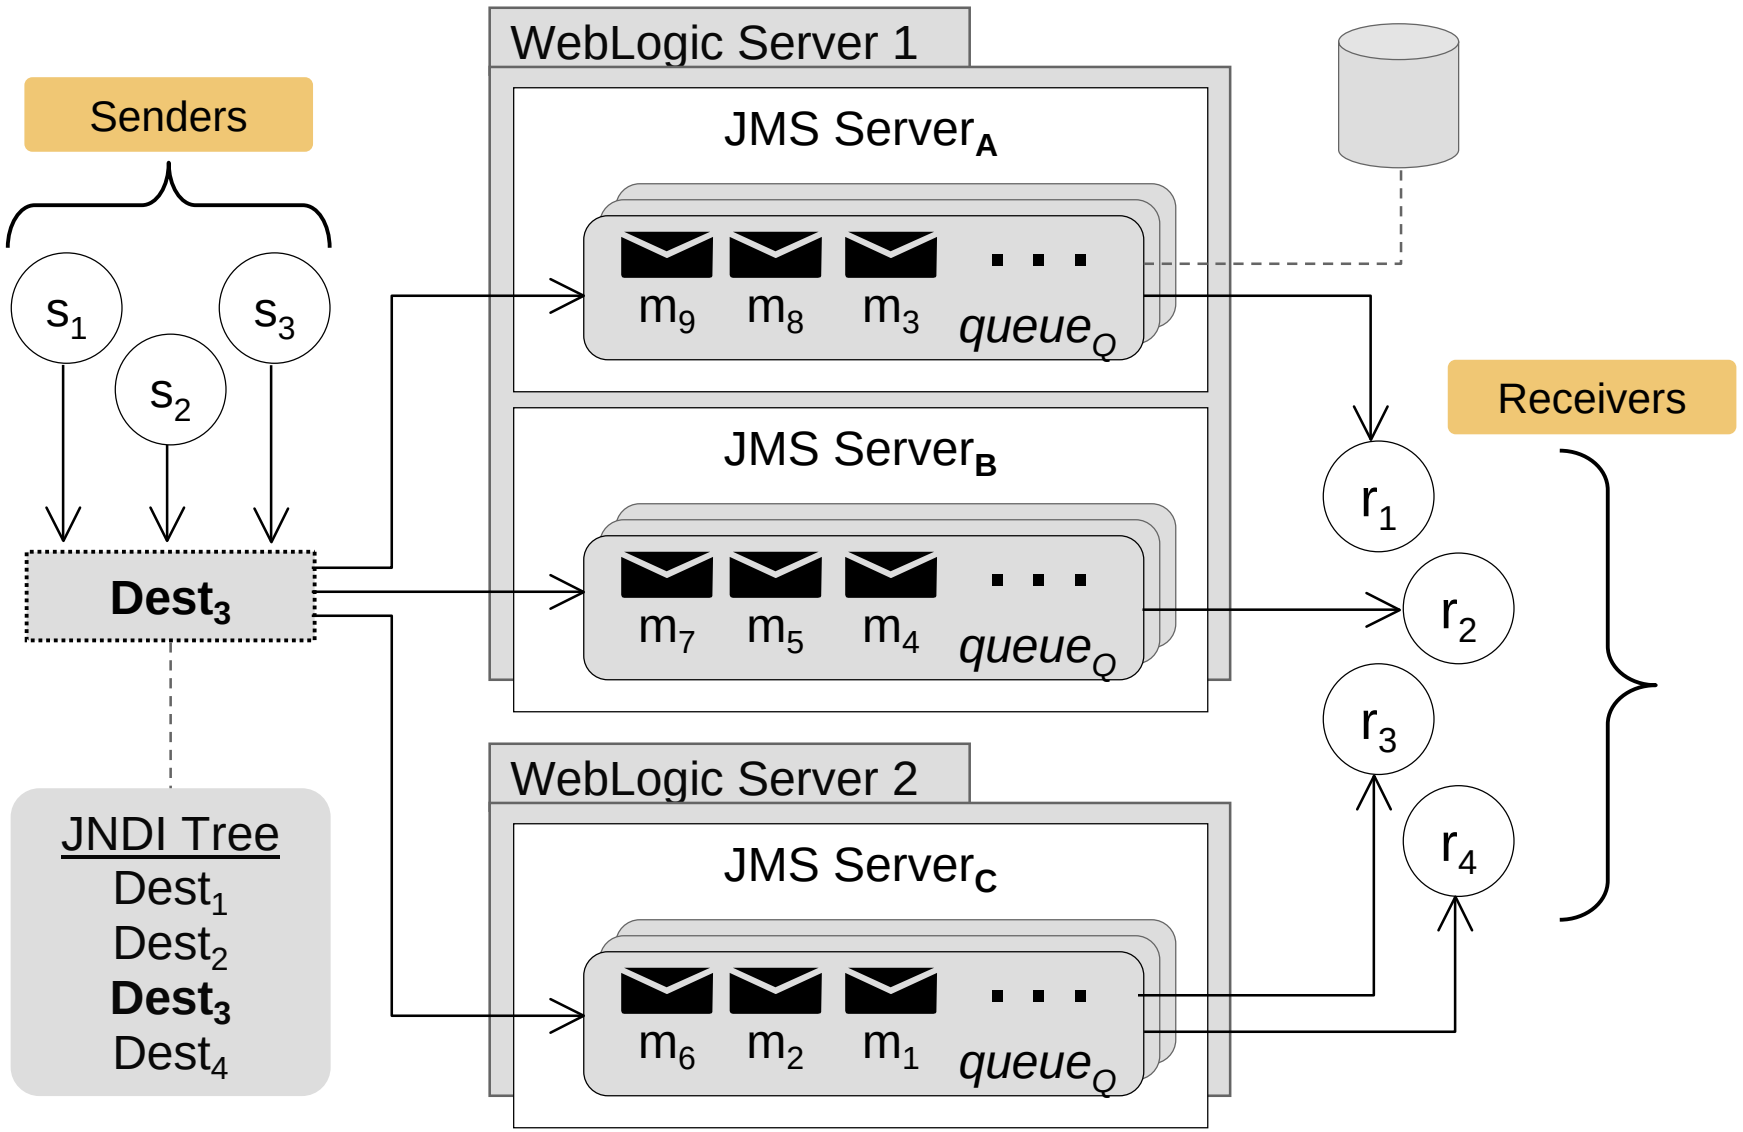

Distributed Topics

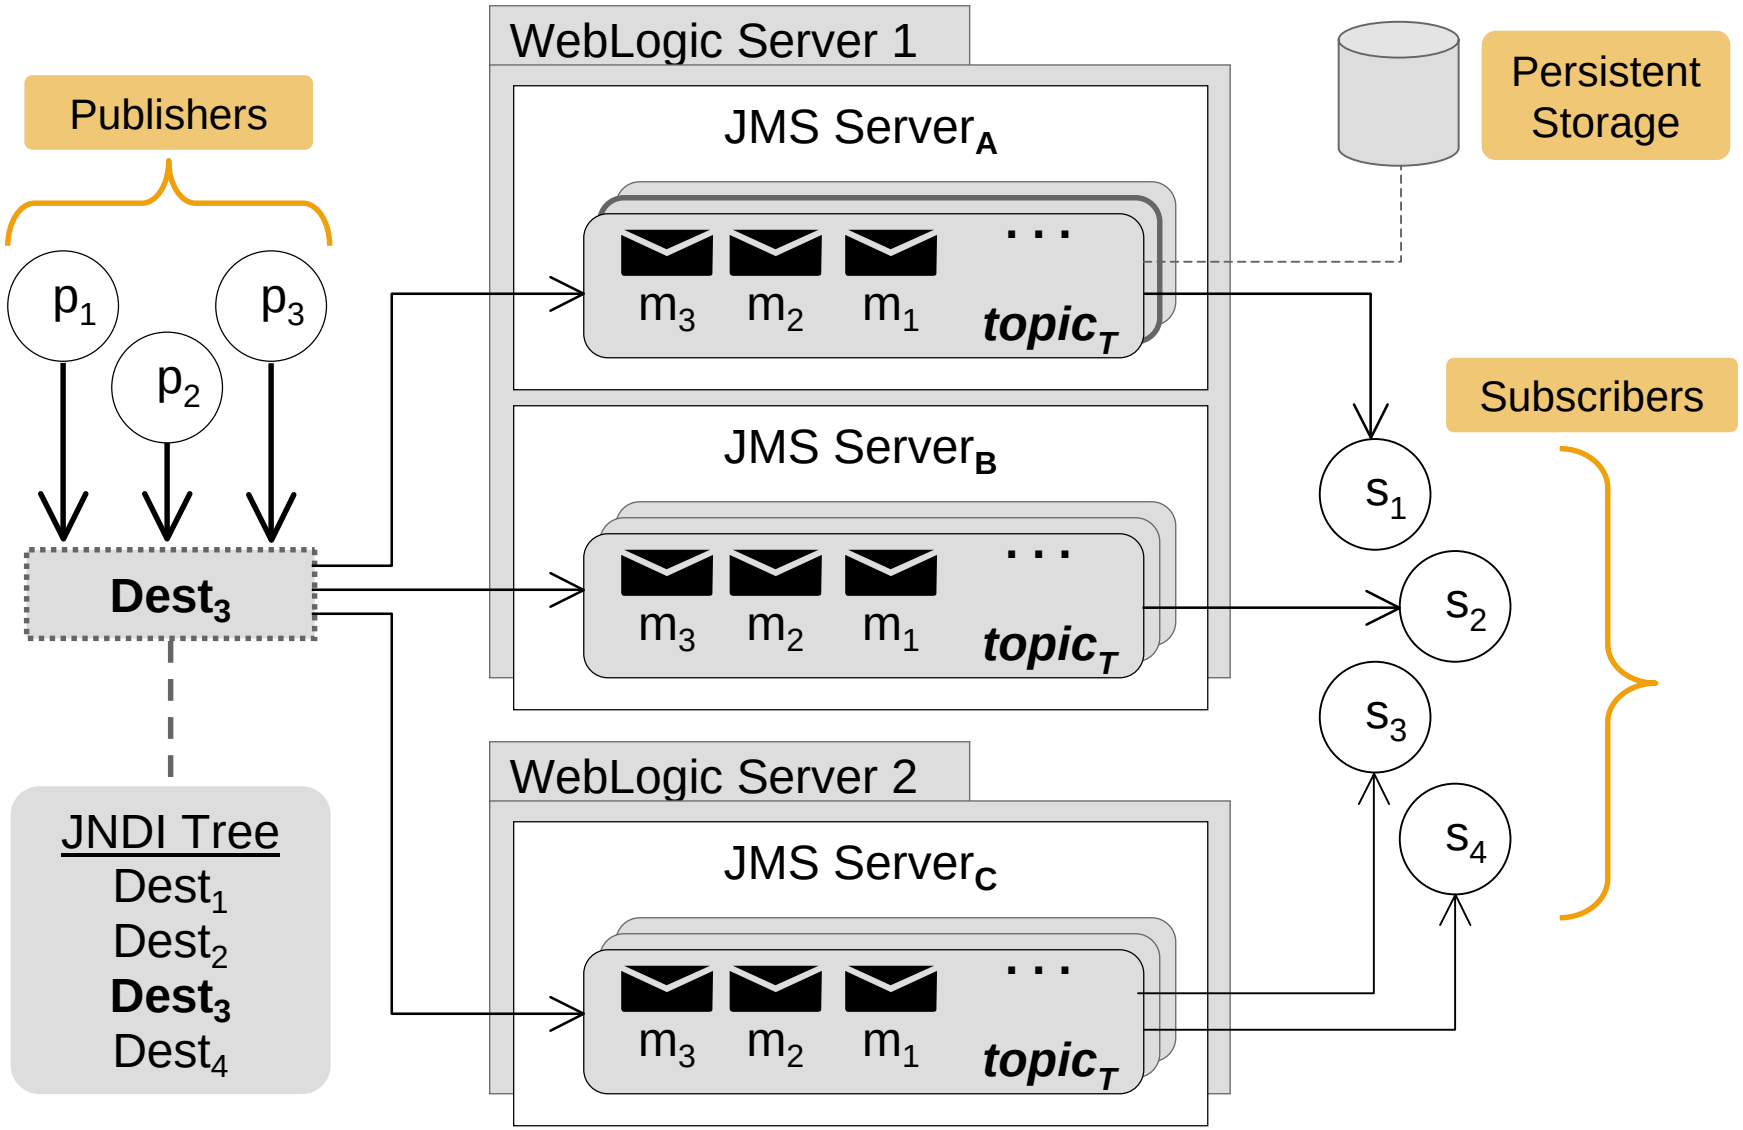

Create a Distributed Topic…

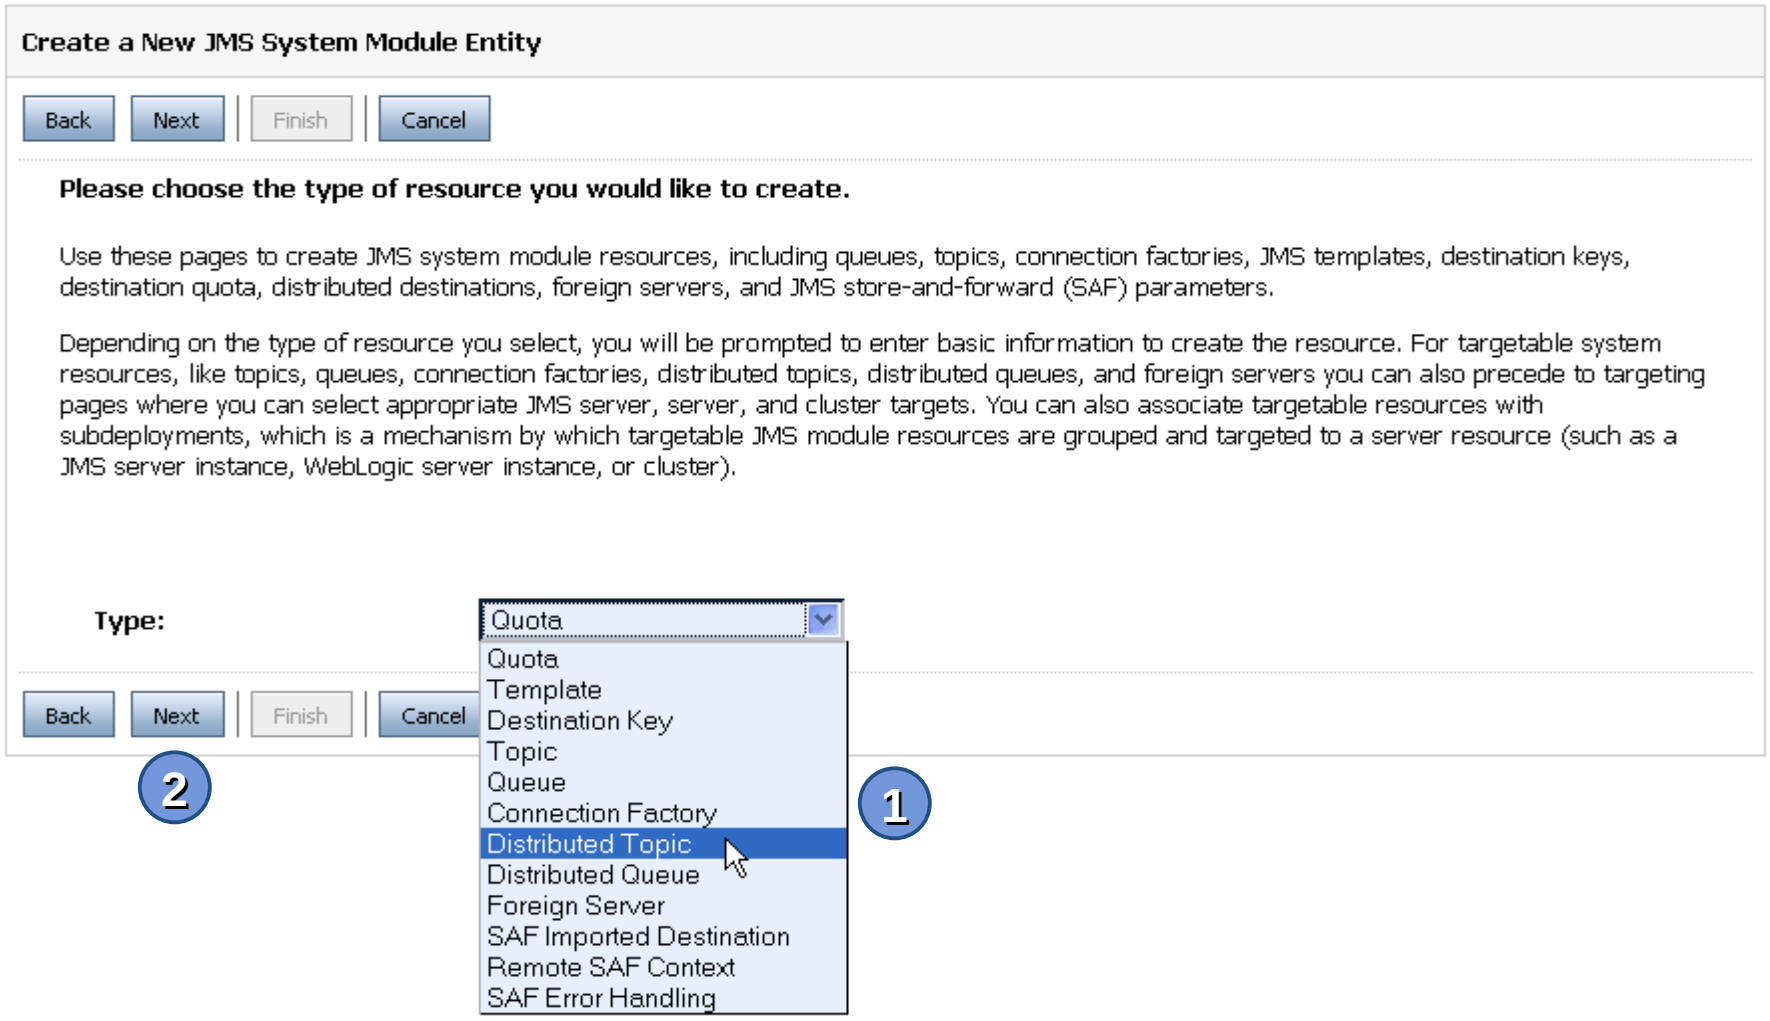

…Create a Distributed Topic…

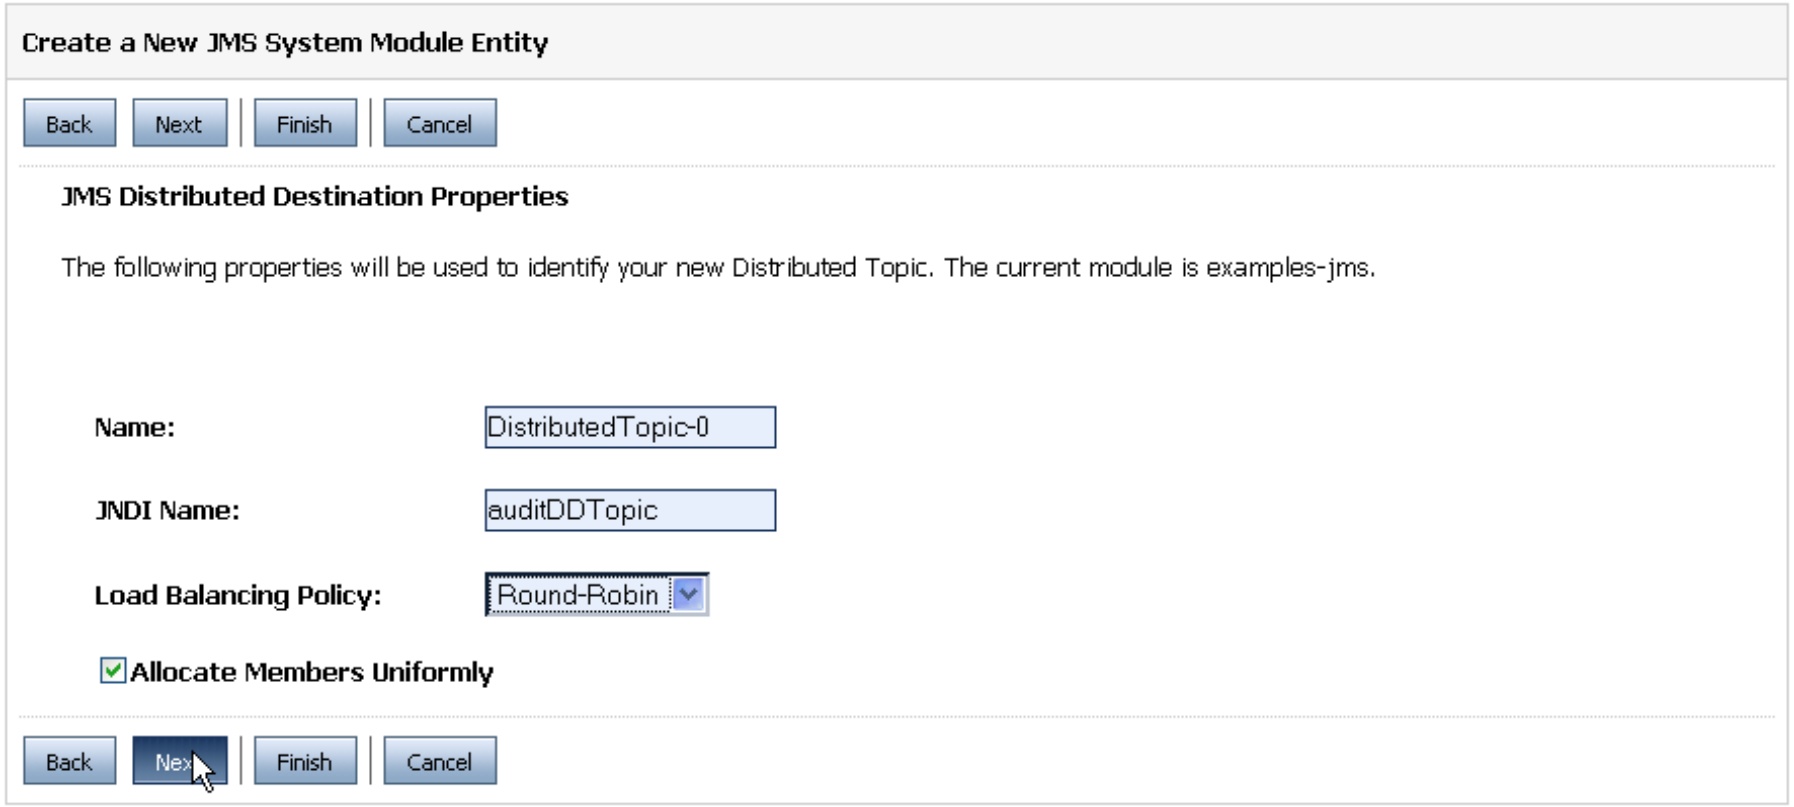

...Create a Distributed Topic

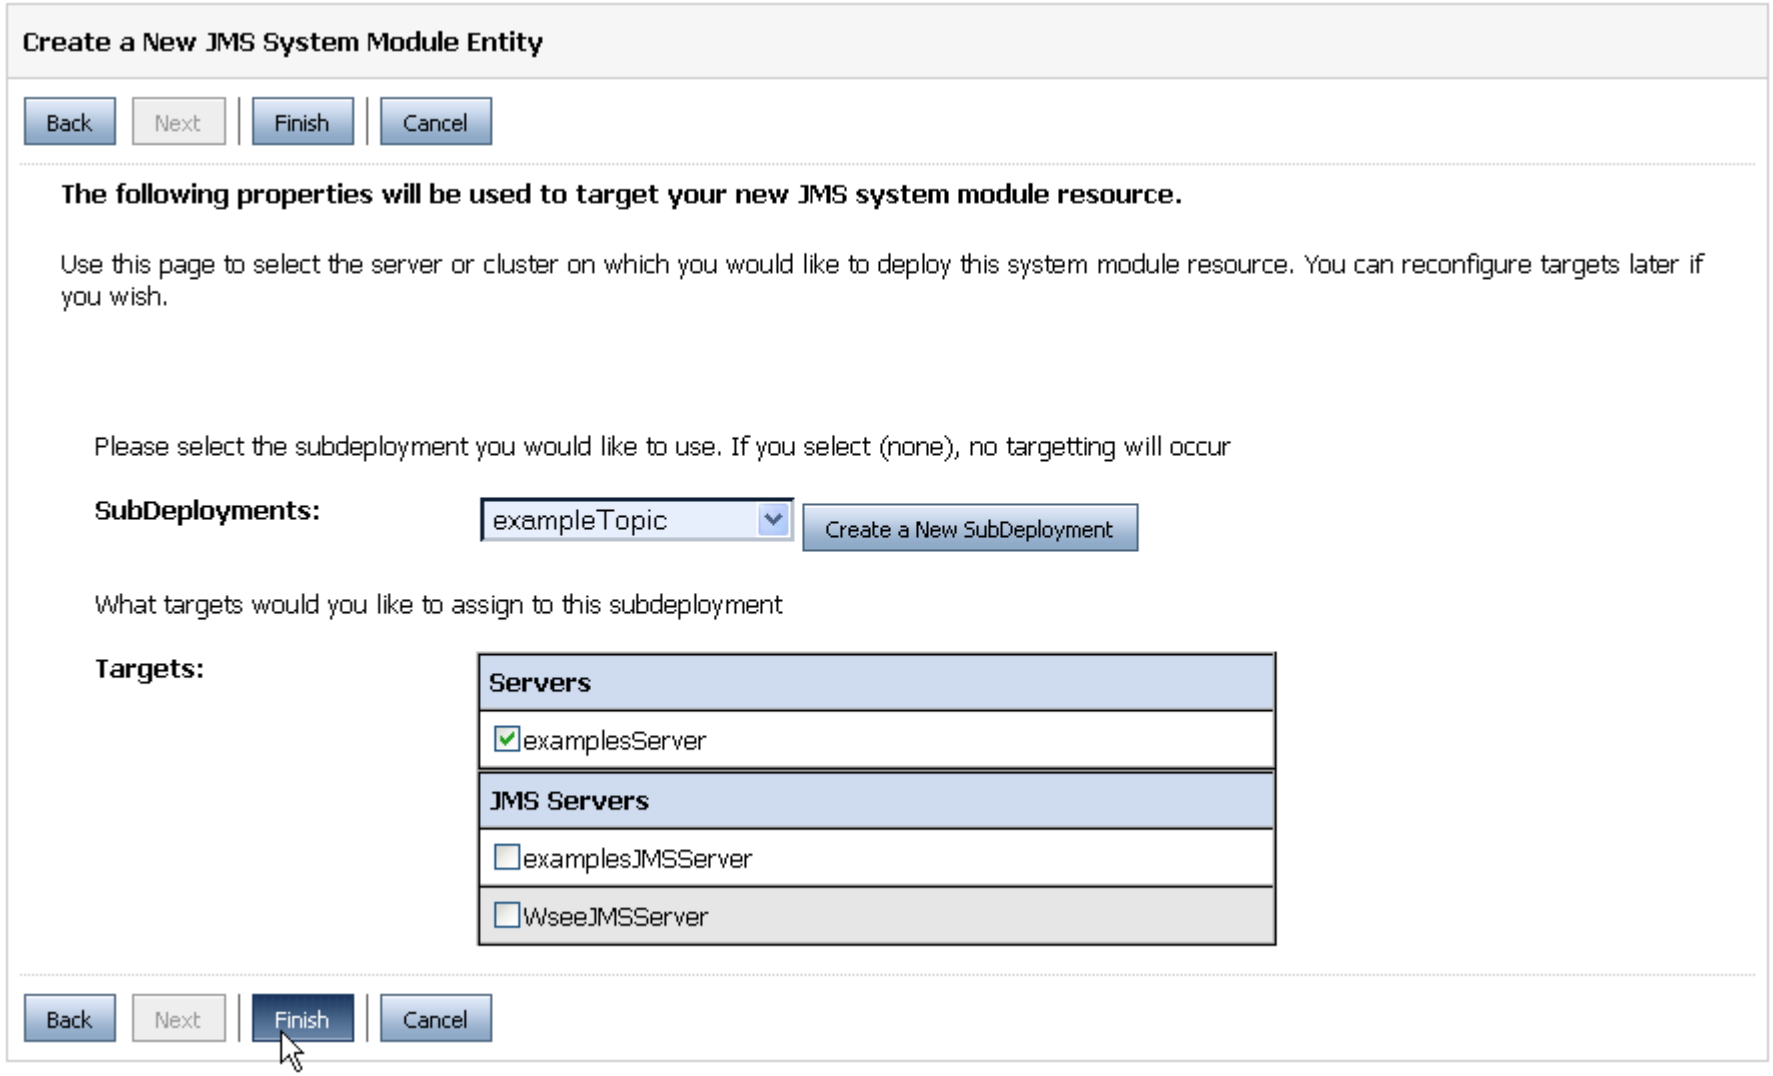

Distributed Destination Threshold and Quota

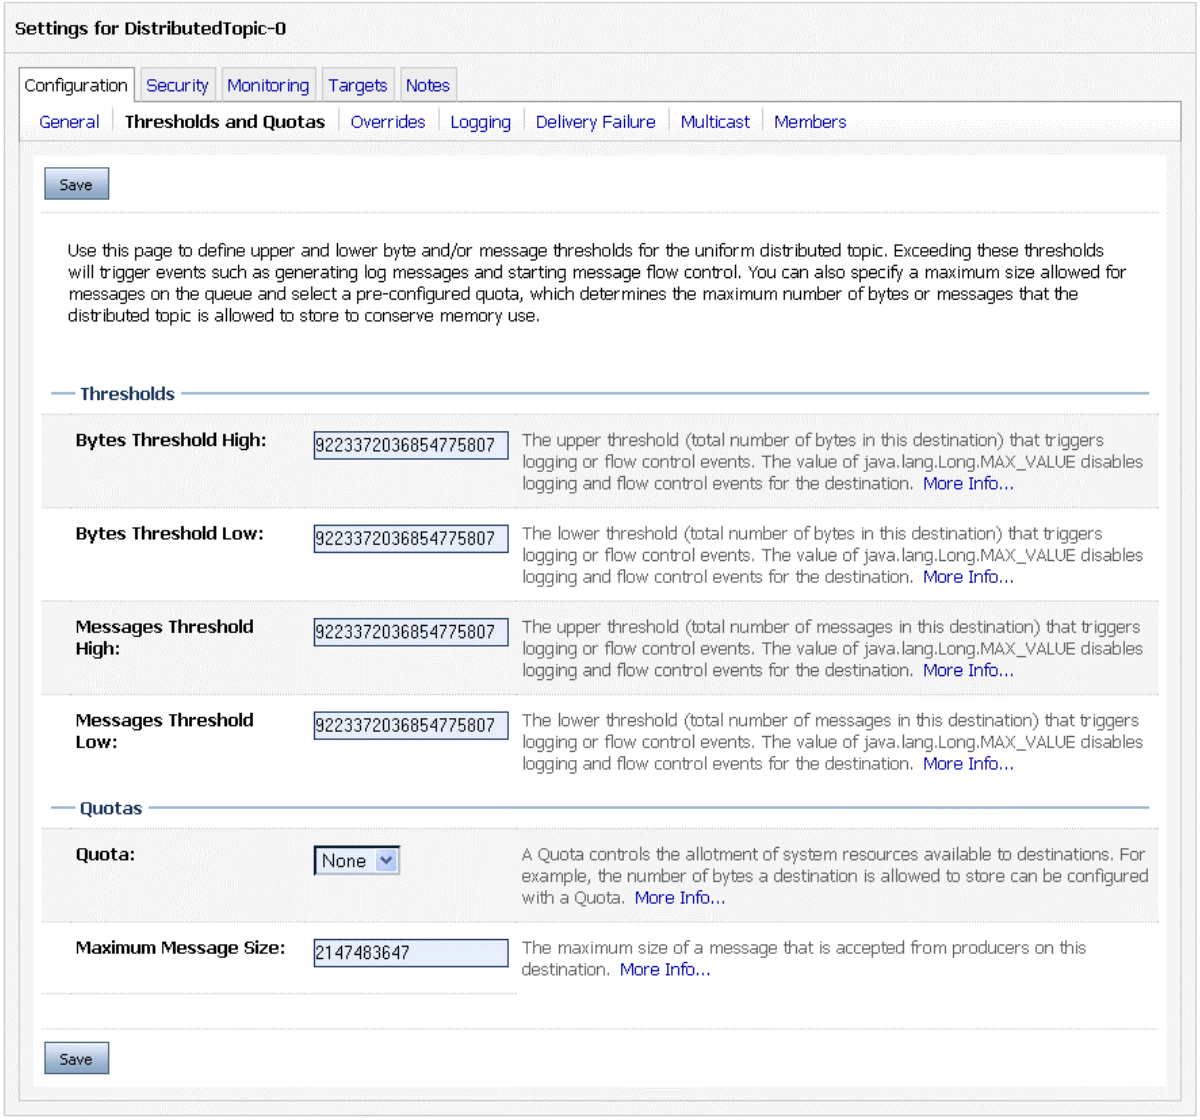

Server Affinity

- WebLogic will try to establish new connections to the same servers as existing connections.
- Applies to distributed destinations only.

For a new JMS connection distributed destination:

1. Attempts to create a connection on the same server as the JNDI initial context, the original client connection server instance.
2. Attempts to create a connection to a server with which the client already has a connection, JMS or non-JMS. This may be JTA or JMS transaction dependent.
3. Creates the connection on any server.

Zero Consumers…

- *Server affinity* favors members on the same server as the client.
- *Transaction affinity* favors members that are already involved in a transaction.
- A *zero consumer* is a member that has no consumers; it is used as follows:
  - A producer avoids a zero consumer
  - A consumer searches a for zero consumer
- *Persistent messages* are put in a backing store.

…Zero Consumers

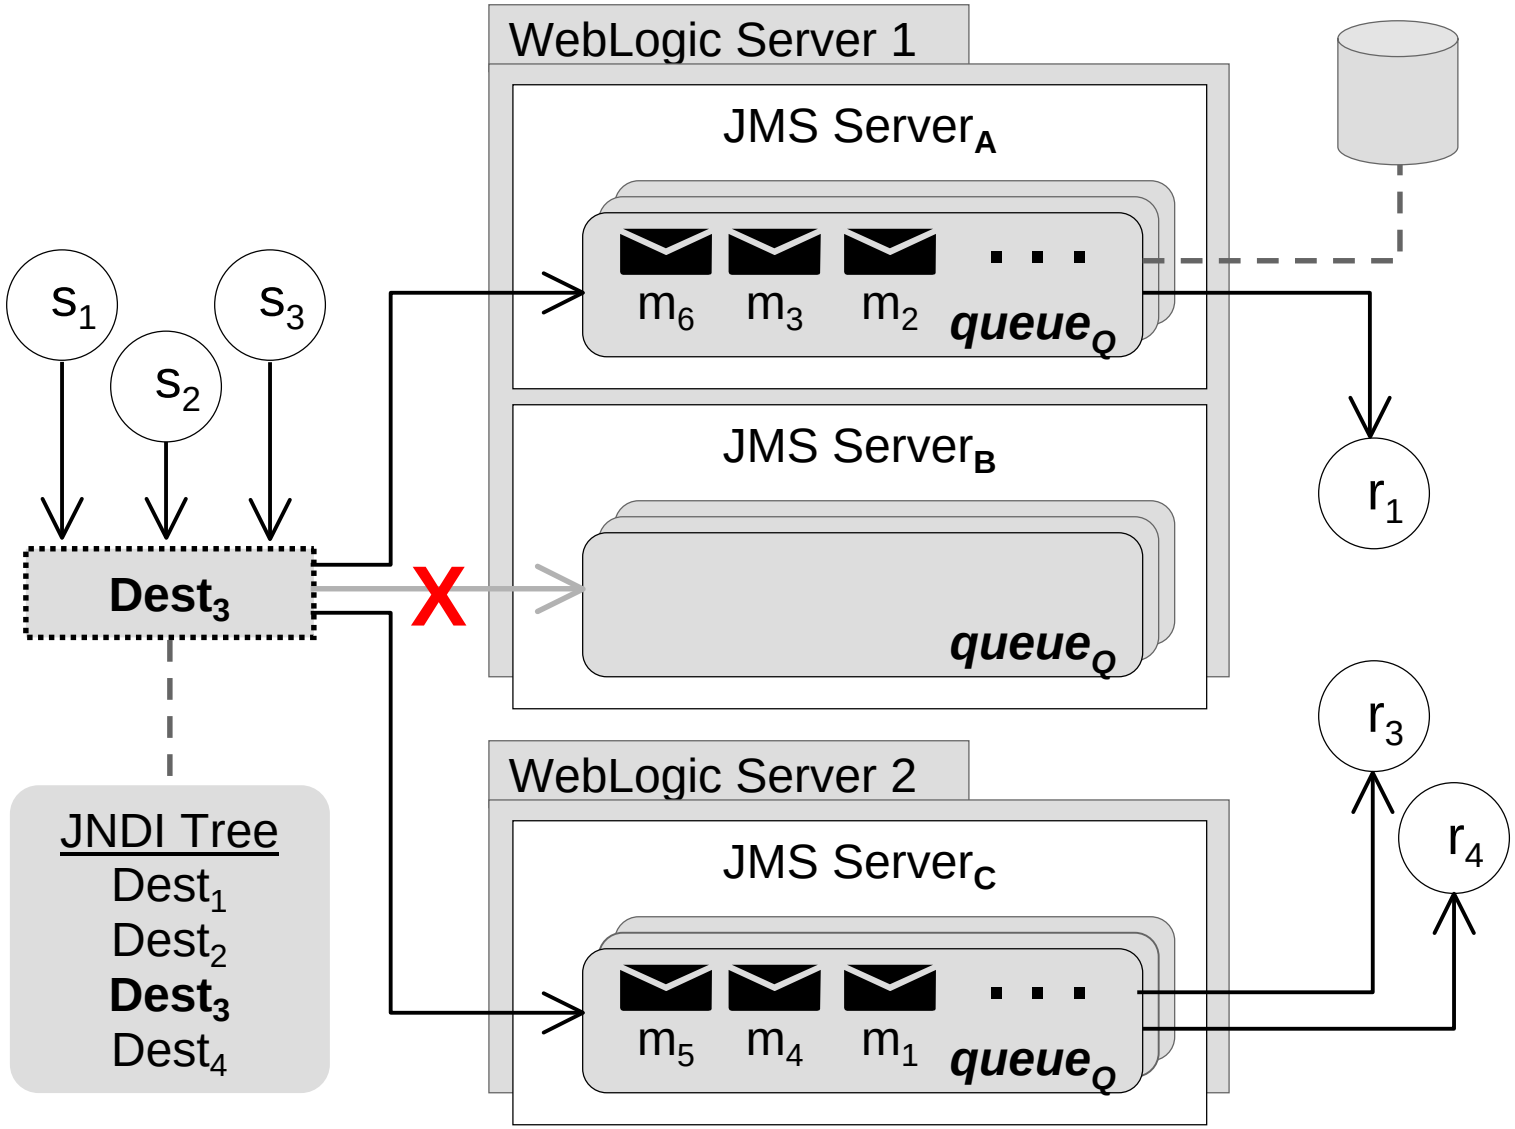

Load Balancing

- If the destination is not distributed:
  - All production is directed at one destination
  - The consumer uses that destination
  - Load balancing does not occur
- If the destination is distributed:
  - Production is replicated across member topics and queues
  - Load balancing is either round-robin or random
  - The consumer is pinned to one member destination, which is selected at creation time

Location Transparency

- Standard destinations can be replicated in the cluster-wide JNDI service.
- Choose *Replicate JNDI Name in Cluster* in the Admin console.
- Use the **`isJNDINameReplicated()`** method in the **`JMSQueueMBean`** or **`JMSTopicMBean`**.
- Clients look up a destination on any server and receive a connection on the actual server.

Failover

- Failover is both automatic and manual.
- A manual failover includes migrating these elements:
  - An entire JMS Server
  - A persistent store
  - A transaction log
- WLS provides a migratable JMS service that attempts to deliver outstanding JMS messages.

JMS Server Migration

- JMS is an *exactly-once* service; each JMS server exists on exactly one WebLogic server.
- When a WebLogic server fails, its JMS servers can be migrated. Migration must be configured ahead of time.
- For persistent messaging, migratable targets must have access to the same JMS store as the original server.
- A JMS server can migrate to a WebLogic server that already hosts distributed destination members.
- Migration may be a part of scheduled maintenance.

Performing Migration

- Services migrated to a non-running server will be started when the server starts.
- Migration must can manually initiated or automatically initiated.
  - Automatic migration happens through whole server migration only
- It is possible to use third-party products to perform migration:
  - JMX
  - Veritas HA

JMS Migratable Targets

- The JMS server can be manually migrated to any other server in a cluster
  - Optionally, an administrator can limit which servers will receive the migrated JMS server by specifying a defined set of migratable targets within a cluster.

Migrate JMS Data

- For the JDBC store:
  - If it is on the failed server, migrate the database to the new server and change the JDBC URL for the JDBC store's DataSource.
  - If it is not on the failed server, no changes are required.
- For the file store:
  - Migrate to a new server
  - Ensure that the pathname is the same on the new server as the original one
- For transactions, also migrate the transaction logs

Migration Configuration

- To configure a JMS migration:
  1. Define migratable target servers in the cluster.
  2. Create JMS Servers and assign to the migratable targets.
  3. Create *physical* member destinations (use the auto-deploy option for distributed destinations).

Configure a Migratable Target

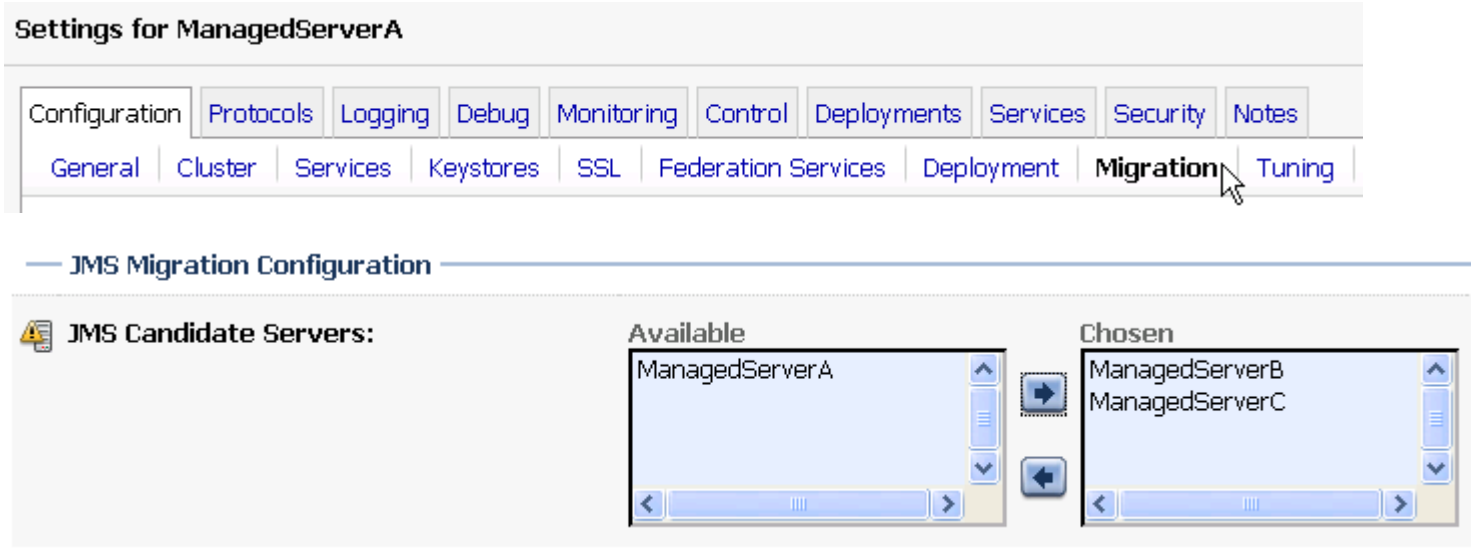

Target JMS Server to a Migratable Target

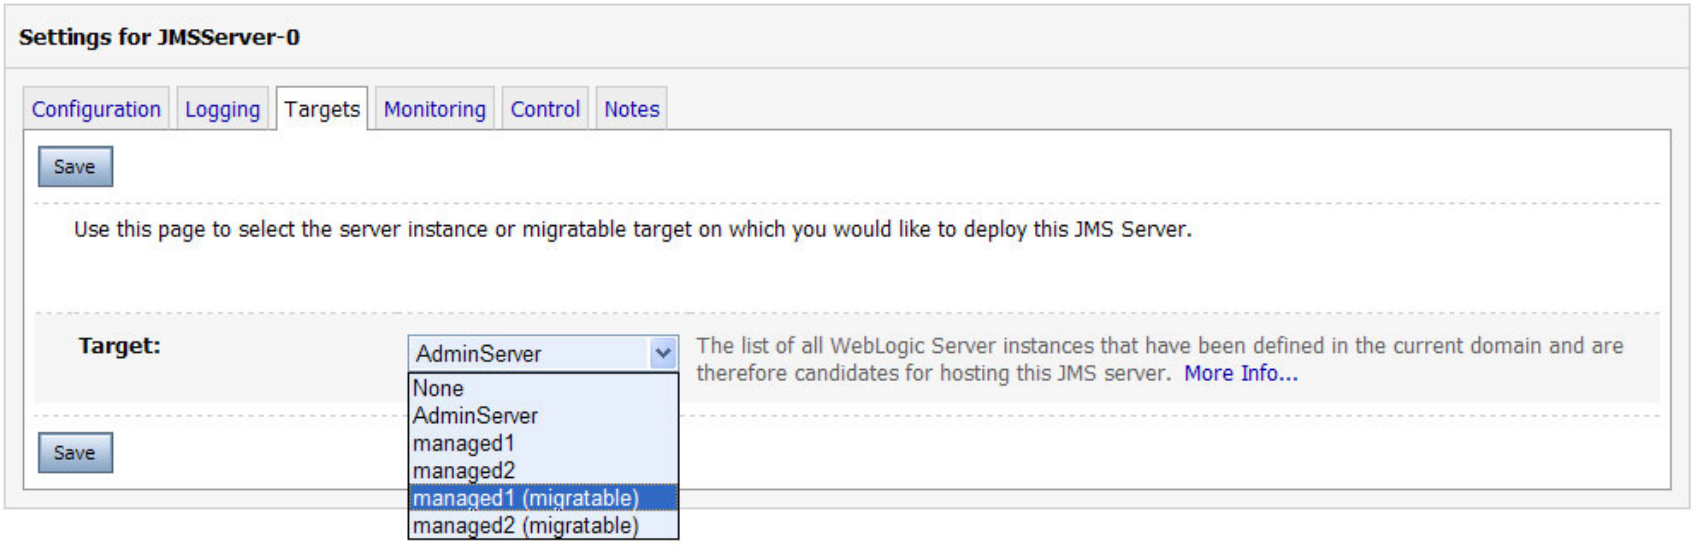

Migrate Services

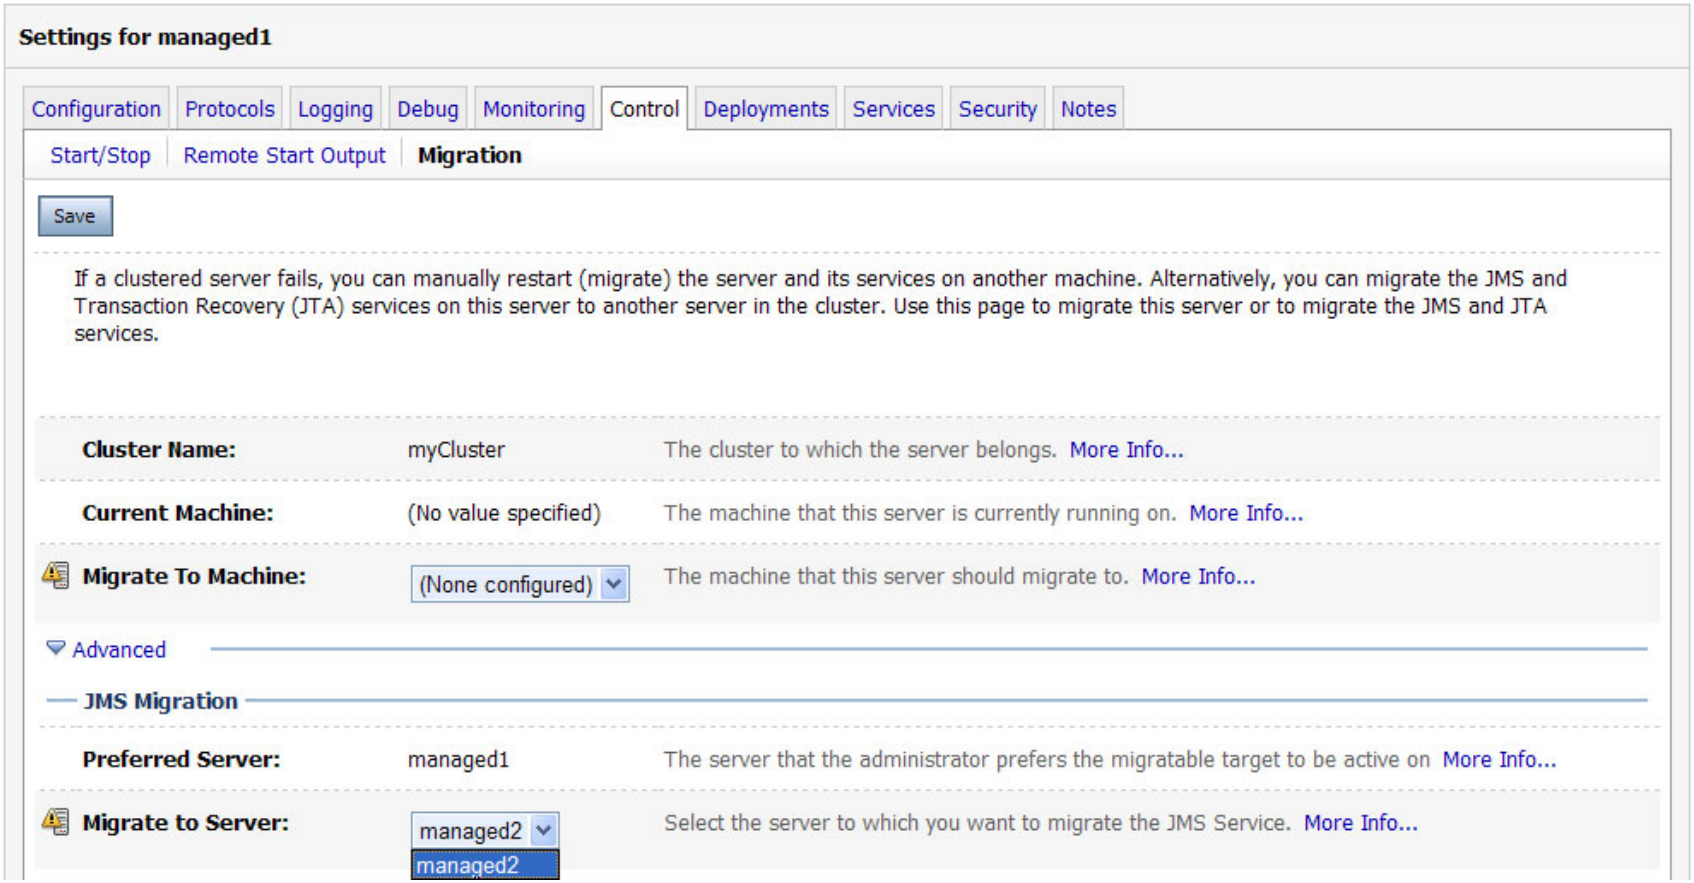

Section Review

In this section, we learned how to:

- ✓ Handle failover and load balancing of JMS connections
- ✓ Cluster a JMS connection factory

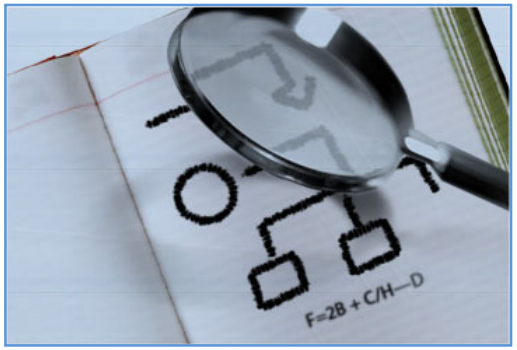

Exercise

Load Balance JMS Messages

- For details on the exercise, refer to the Lab Guide.
- If questions arise, ask the instructor.
- The instructor will determine the stop time.

Module Review

In this module, we learned how to:

- ✓ Migrate a whole server
- ✓ Cluster JNDI objects
- ✓ Cluster JDBC data sources
- ✓ Cluster transactions
- ✓ Cluster JMS connection factories